

# Neural Networks and Deep Learning

## Lecture 1: introduction

Jan Chorowski  
Instytut Informatyki  
Wydział Matematyki i Informatyki Uniwersytet  
Wrocławski  
2017-2018

# Practical Information

- Website:  
<http://ii.uni.wroc.pl/~jch/teaching/neural-networks-fall-17>
- Free on-line Textbooks:
  - <http://www.deeplearningbook.org/> -> You can also buy a printed version!
  - <http://neuralnetworksanddeeplearning.com/>
- Good lecture notes for 50% of material in this course:  
<http://cs229.stanford.edu/>
- More advanced, very good courses on deep learning:
  - <https://ift6266h15.wordpress.com/>
  - <http://cs231n.stanford.edu/>
  - <http://cs224d.stanford.edu/>
- Good paper textbooks:
  - Goodfellow, Bengio, Courville „Deep Learning”  
<http://www.deeplearningbook.org/>
  - Murphy, Machine Learning: a Probabilistic Perspective  
<http://www.cs.ubc.ca/~murphyk/MLbook/>
  - Bishop, Pattern Recognition and Machine Learning (PRML)  
<https://www.microsoft.com/en-us/research/people/cmbishop>

# Why are you here?

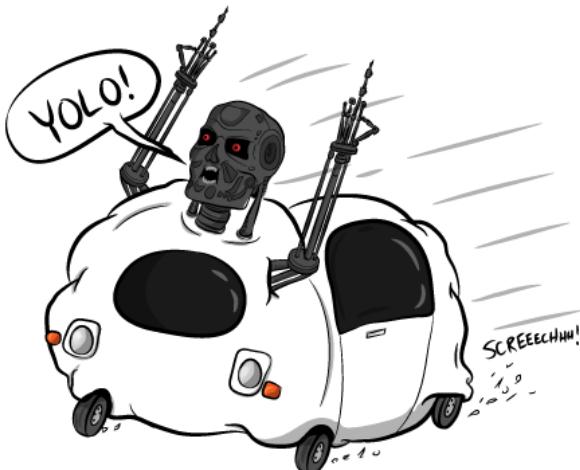
What do you want to learn?

1. Machine Learning: using data to teach machines useful skills
2. Neural Networks: good Machine Learning models

# Where is machine learning?

## Everywhere!

- web search ☺
- and ads ☹
- recommendations
- Self-driving cars



Local results for **starbucks** near **Chicago, IL**

**Ads**

**Starbucks** Get Local Directi...  
Phone Numbers |  
MapQuest.com

**Starbucks Ct** Whatever you're l...  
you can get it on  
www.eBay.com

**Buy Starbuck** Find Starbucks C...  
eBay Express Of...  
www.eBayExpres...

**Local Search Results**

**Starbucks in Chicagoland**  
This friendly neighborhood Starbucks is extra-spacious, ... of local hero Joe DiMaggio in this first Starbucks in the Little Italy neighborhood of Chicago. ...  
[www.starbuckseverywhere.net/Chicagoland.htm](http://www.starbuckseverywhere.net/Chicagoland.htm) - 127k -  
Cached - Similar pages

**Starbucks in Illinois**  
Illinois Chicagoland - Illinois Remote  
[www.starbuckseverywhere.net/illinois](http://www.starbuckseverywhere.net/illinois)  
Cached - Similar pages

**Google Organic Search Results**

amazon.com

Hello, Scott Wheeler. We have recommendations for you. (Not Scott?)  
Scott's Amazon.com Today's Deals Gifts & Wish Lists Gift Cards

Shop All Departments Search Amazon.com

Scott, Welcome to Your Amazon.com (If you're not Scott Wheeler, click here.)

**Today's Recommendations For You**

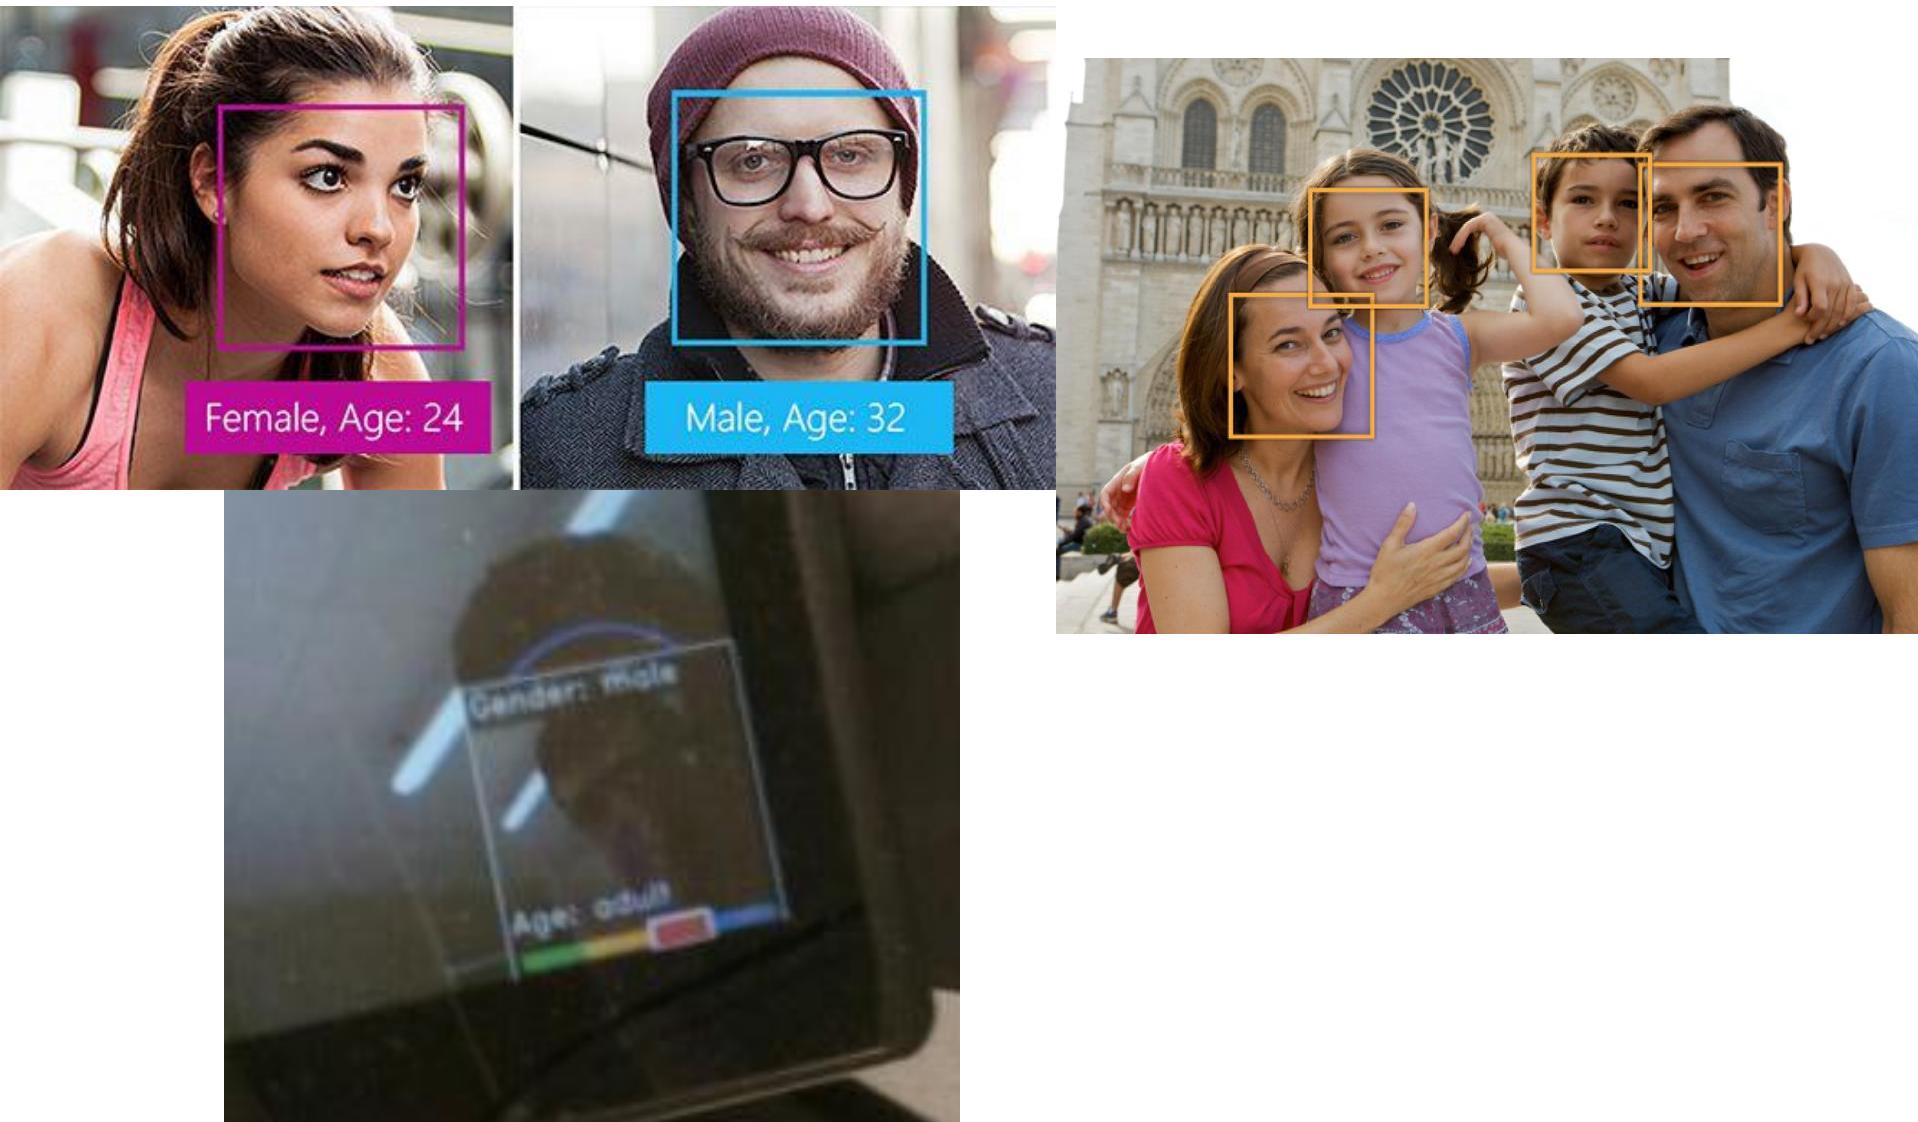
Here's a daily sample of items recommended for you. Click here to [see all recommendations](#).

**Russia Map** by ITMB International Travel Maps By ITMB Publishing Ltd \$10.95

**In Search of Sunrise, Vol. 2: Asia** by Ol' Theta ShinkinBank! \$13.99

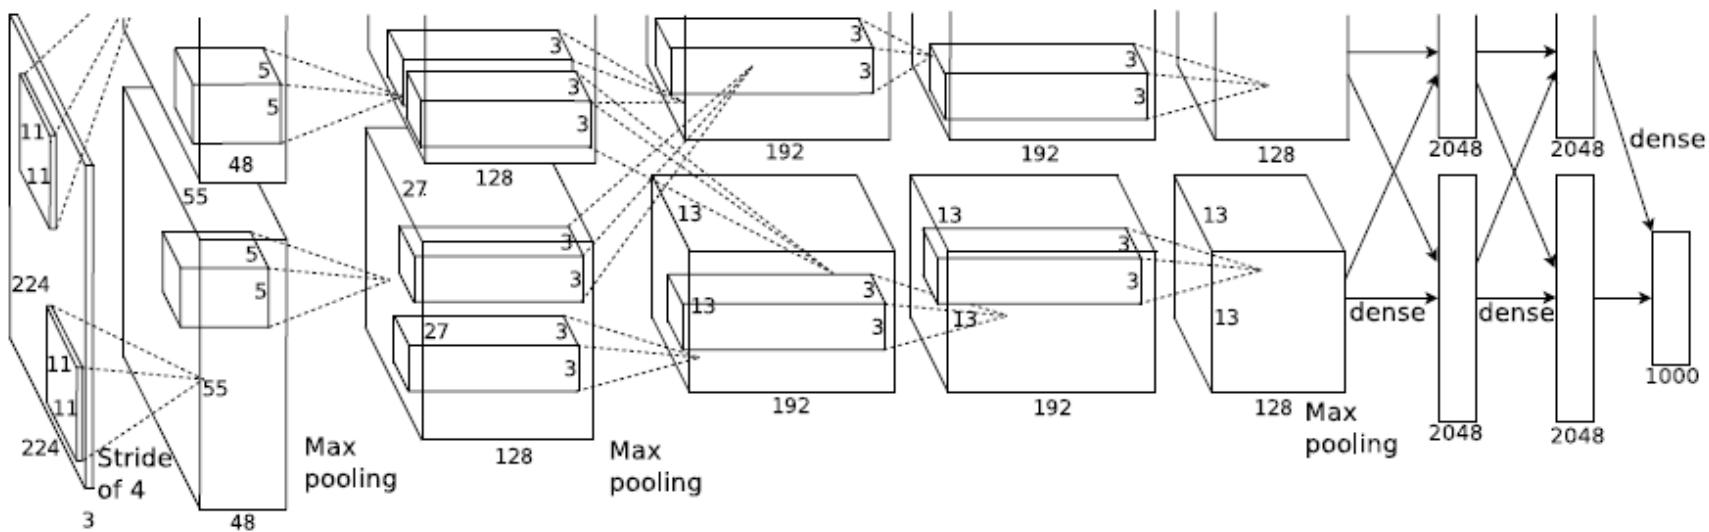
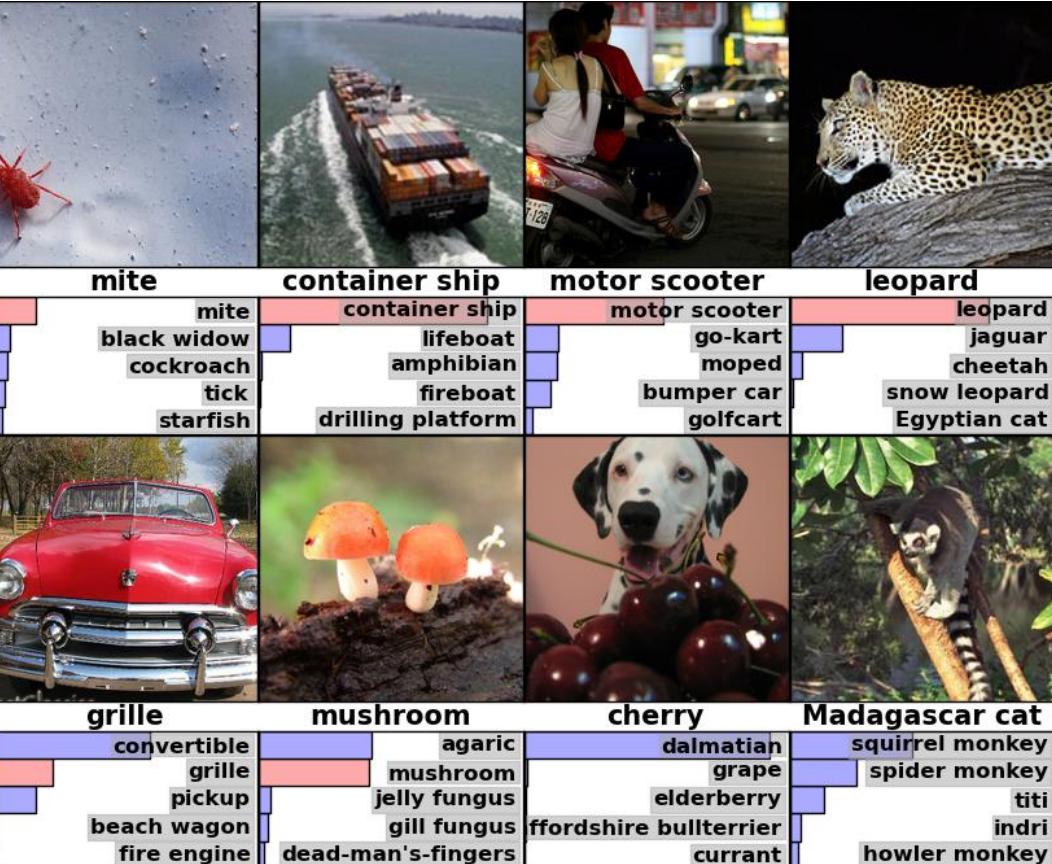
**Land of the Horizons: A History of the Oba** by Jason Goodwin \$11.95

# ML in action: face detection



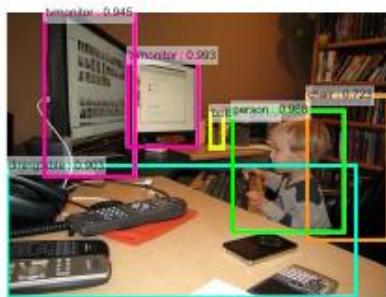
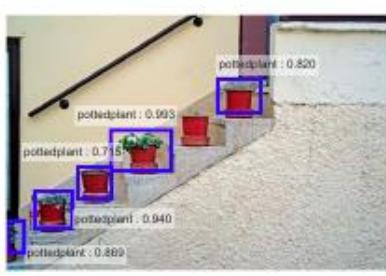
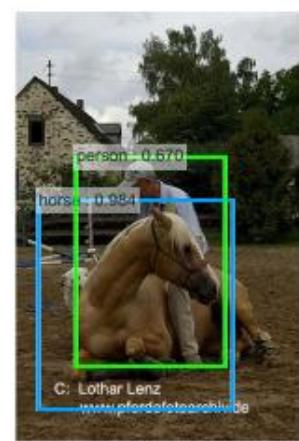
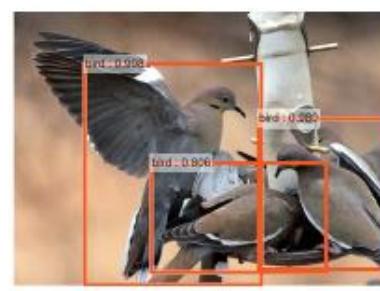
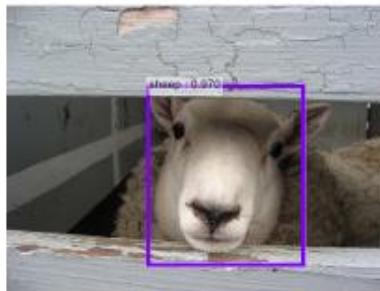
Source: Microsoft and Apple face detection API documentations, wykop.pl

# Image recognition



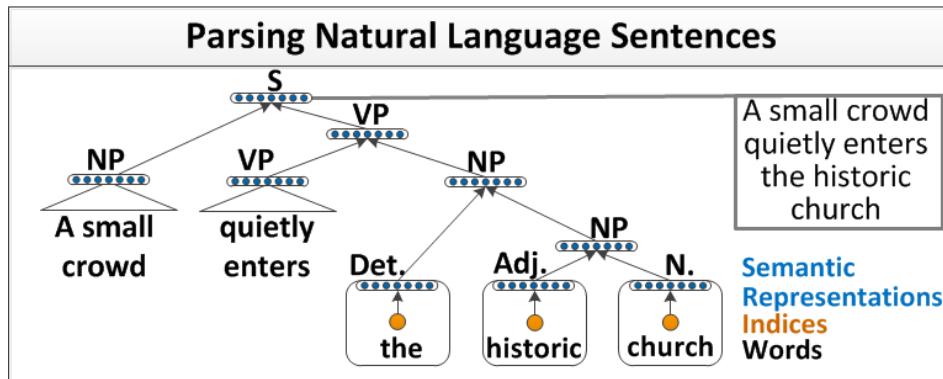
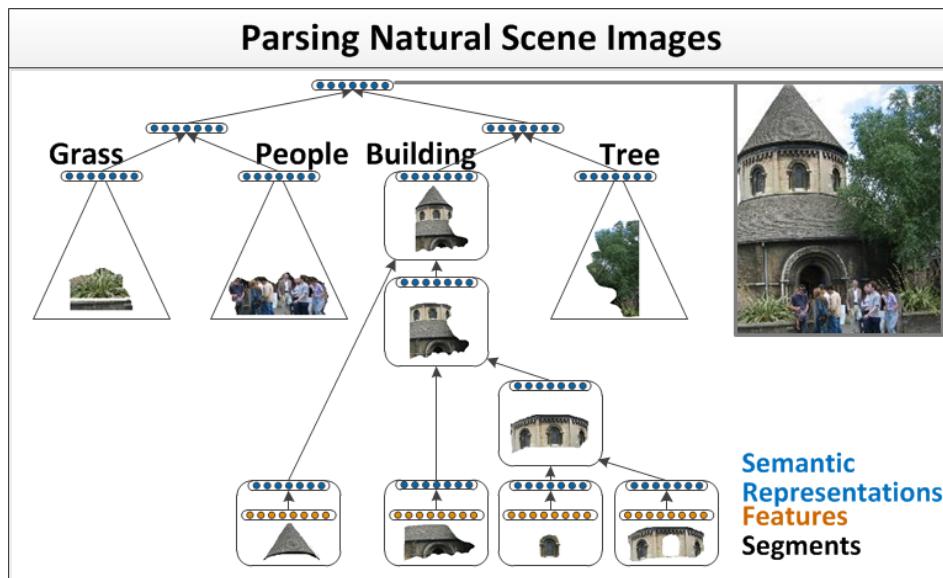
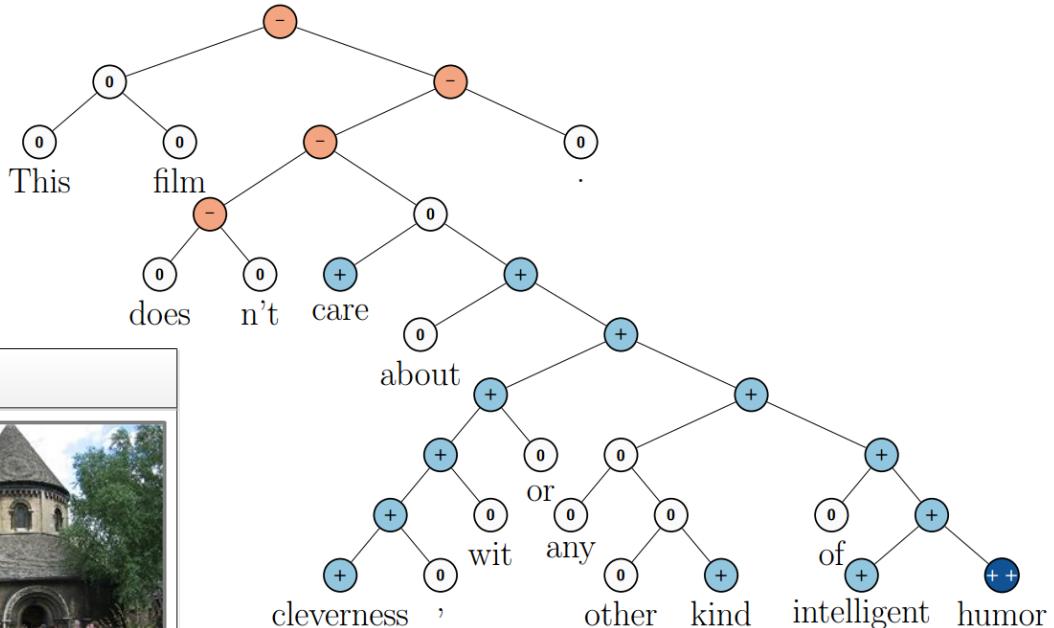
# Object Detection

<https://www.youtube.com/watch?v=WZmSMkK9VuA>



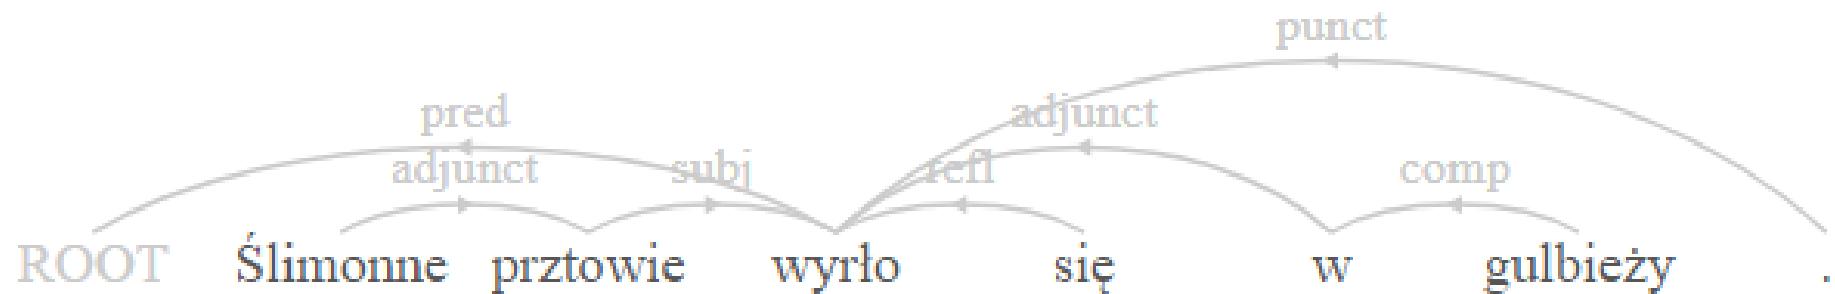
(Ren, He, Girshick and Sun, „Faster R-CNN”, 2015)

# Natural Language Processing (NLP)



# NLP for Polish

<http://zapotoczny.pl/parser/>



# Using both language and images

Putting images and tags into the same vector space  
(Kiros, Salakhutdinov, Zemel, TACL 2015)



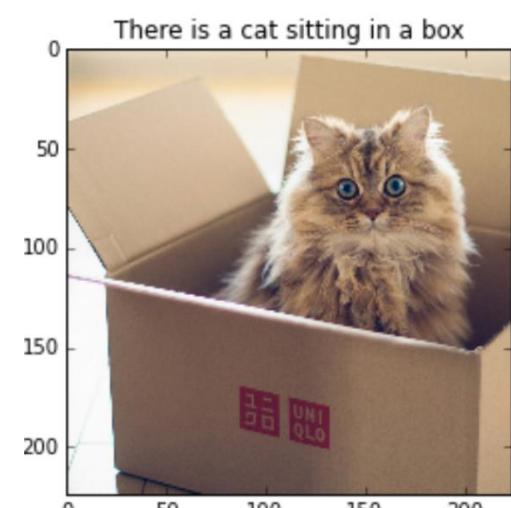
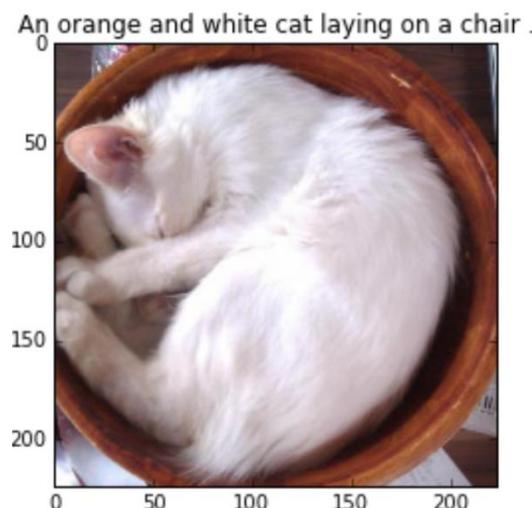
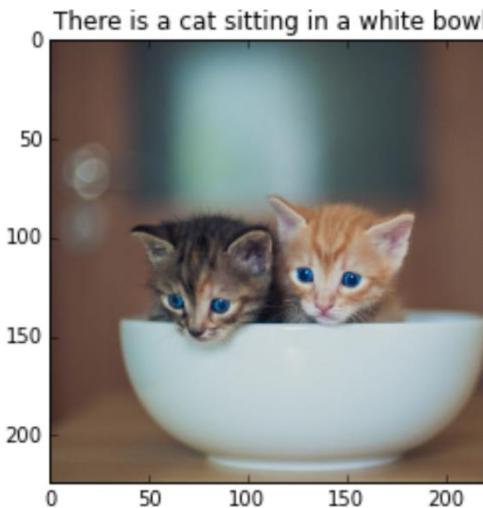
- bowl + box =



- box + bowl =



Caption generation (Xu et al ICML 2015)

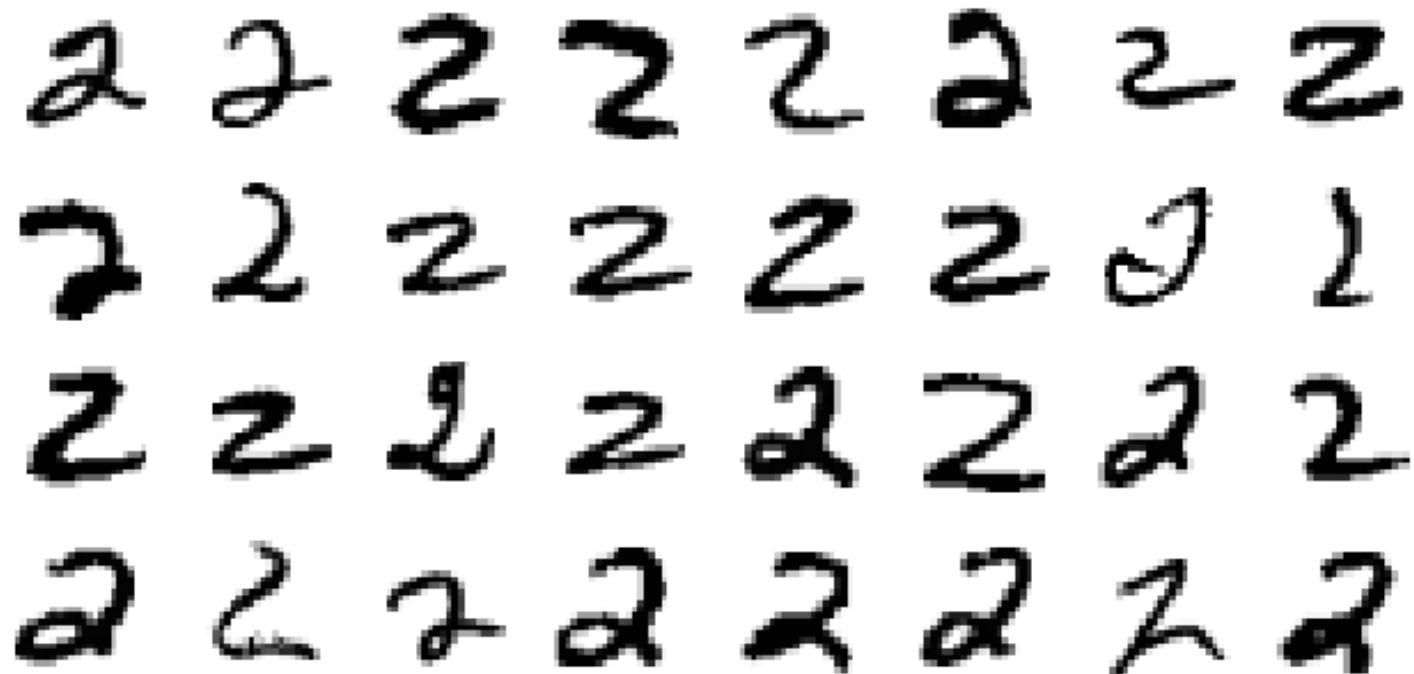


# How are these problems solved?

1. Take **data** and a **learning algorithm**
2. The algorithm discovers **patterns** in the data and produces a **model**
3. Query model to ask questions about the data
  - „Is there any face in the image?”
  - “Is this review favorable?”
  - „What is the object in the image?”
  - “What caption is most likely given this image?”

# Example: Digit Recognition

- Task: recognize handwritten digits
- Input: images  $28 \times 28$  pixel values ( $[0,1]^{784}$ )
- Output:  $\{0,1,\dots,9\}$



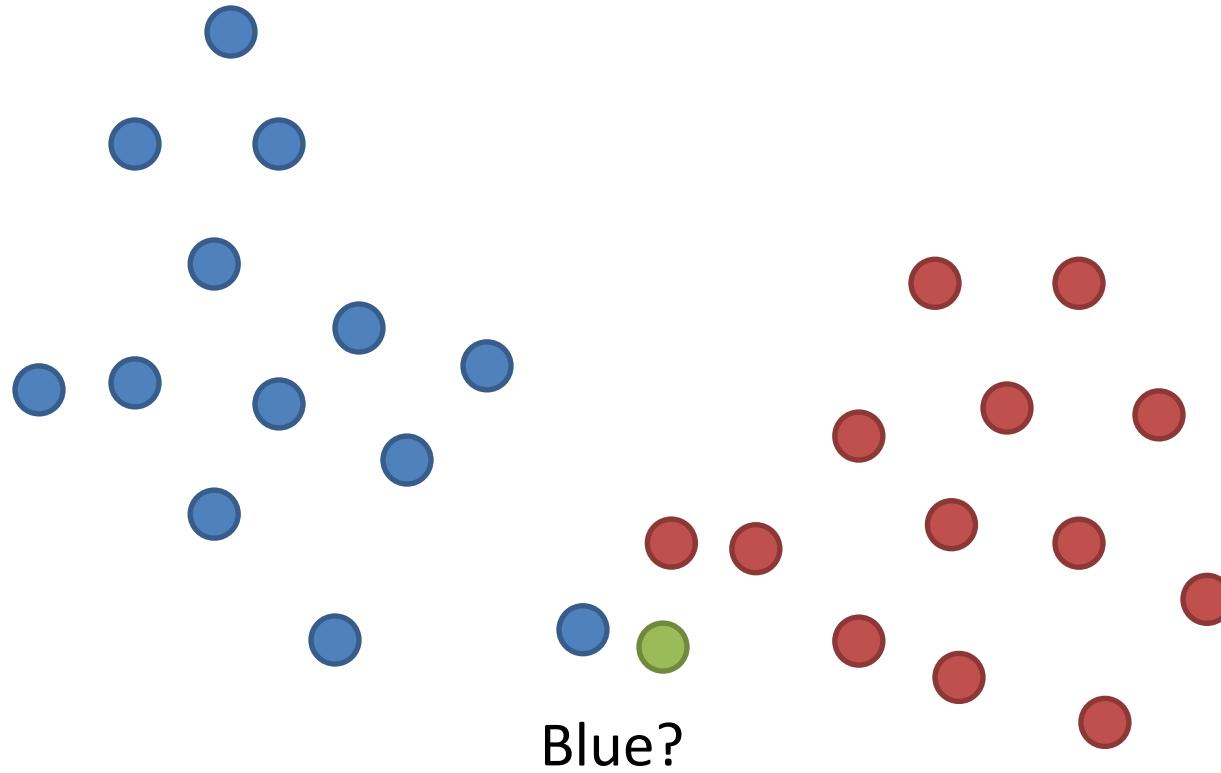
# The easiest algorithm

- Collect examples – pairs of (input, output):

5	0	4	1	9	2	1	3	1	4
5	0	4	1	9	2	1	3	1	4

- To classify a new instance 3:
  - Find the most similar element of the train set 3
  - Return its label
- Seems too easy
  - It works well when we have enough data

# KNN (k Nearest Neighbors)



Design decision: choose “ $k$ ” – how many neighbors to use?  
Few – sensitive to outliers, Many- very smooth classification boundary

# KNN uses: recommendations

- Items are similar if users rate them similarly
- Training (offline):
  - Represent each item by its scores
  - Compute the distance between scores
- Recommendation generation (online):
  - Find nearest neighbors for each item recently browsed/put into the cart

Amazon patented this:

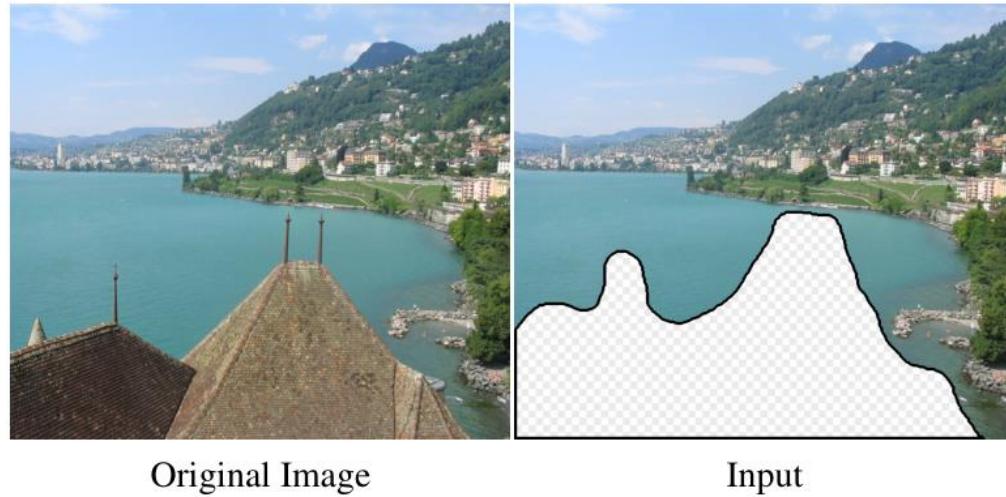


US007113917B2

(12) **United States Patent**  
Jacobi et al.

(10) Patent No.: US 7,113,917 B2  
(45) Date of Patent: \*Sep. 26, 2006

# Nearest Neighbors use: inpainting



Original Image

1. Find similar photos
  - Use Gist descriptors
2. Match and copy

Works poorly with 1000 images...

But is good with  
 $>1000000$



Scene Matches

Output

# Alternative realities



Input



Alternative inpainting results

# kNN for dialogs

A system Adrian Łaćucki coded based on Reddit.

Real (cherry-picked) conversation from Festiwal nauki:

User: hi

Bot: Hi.

User: i like you

Bot: Your in if u want.

User: lmao

Bot: Ayy lmao.

User: ayyy

Bot: Who are you?

User: im a nurse

Bot: Gott im Himmel. You have become the host.

User: and what are you ?

Bot: I am just a talking robot. But I like my job.

User: thats great

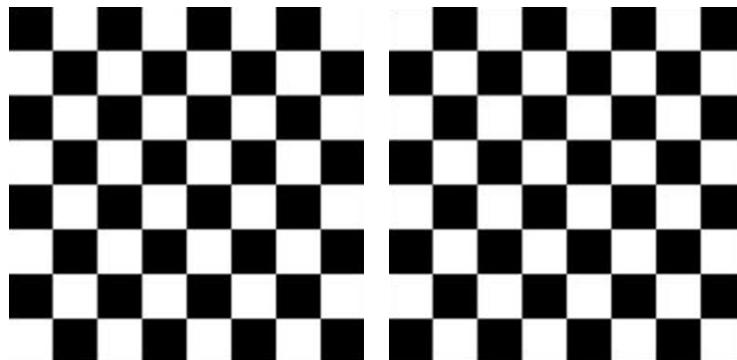
# KNN – what makes images similar?



Quite similar



Not quite similar



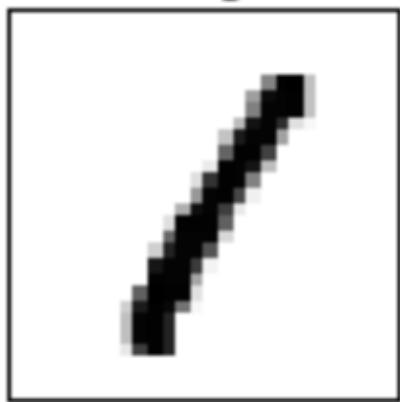
Opposite



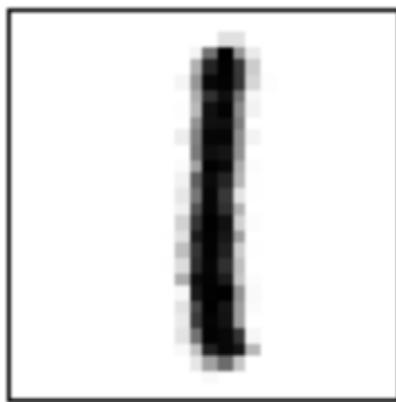
# KNN – what makes images similar?

Pixel-wise difference can be large for visually similar images

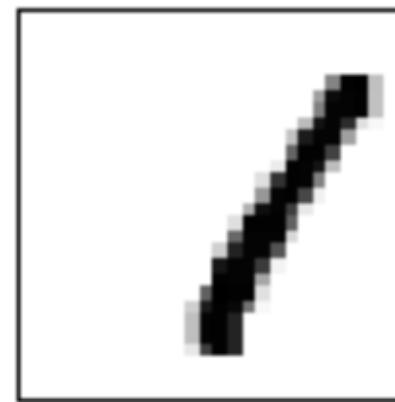
Image



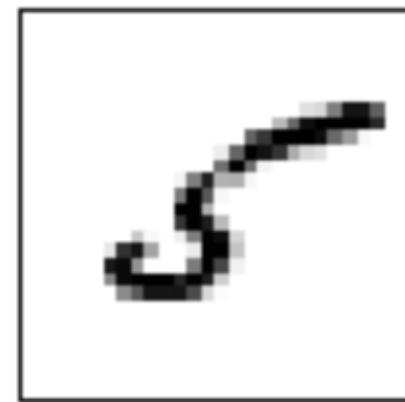
Rotated Image  
diff=76.177262



Translated Image  
diff=114.726490



Other Image  
diff=56.507974



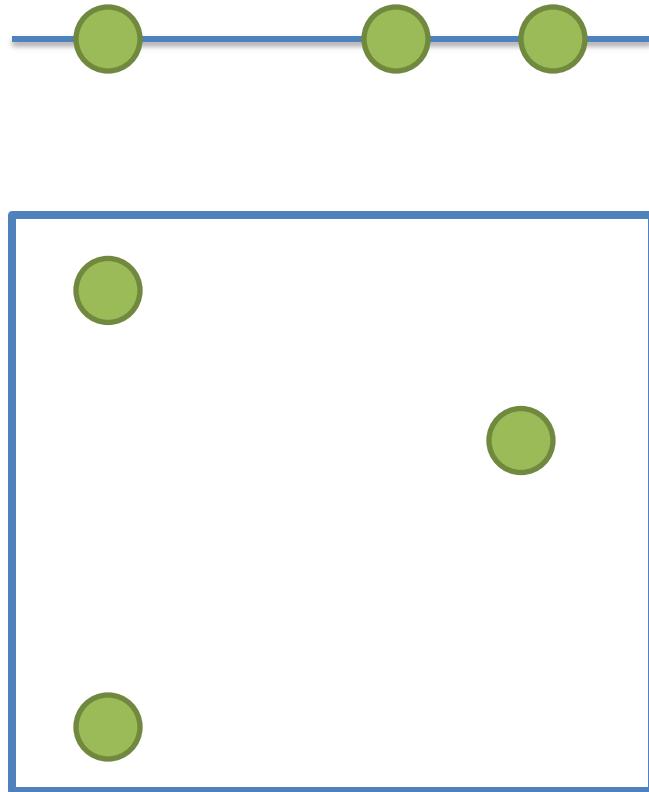
# KNN requires lots of data

- Consider recognizing postal codes (5 digits)
- Will this work if we treat each postal code as a single image ( $10^5$  combinations)?
- If we treat each digit separately, how to accommodate for:
  - Uniform style (slant, thickness) of digits in one code.
  - Segmentation, joins, overlaps of digits.

# Curse of dimensionality

In highly dimensional spaces, things are dissimilar (far away from each other):

- Let  $N$  be the number of points.
- The average distance between points grows with data dimensionality.
- In other words, there are fewer neighbors within a radius from each point.



# Humans see differently

## Change blindness

- Humans are blind even to very drastic changes in images.
- We clearly treat things as the same even if the pixel-wise differences are huge!
- Example:

<https://www.youtube.com/watch?v=1nL5ulsWMYc>

<https://www.youtube.com/watch?v=FWSxSQsspiQ>

# Learning is about invariants and generalizations



# We want a hierarchical model



Low-level  
concepts:

- Edges
- Points
- Gradients
- ....



Higher-level:

- Textures:
  - Spots
  - Stripes
- Geom.  
shapes:
  - Corners
  - Circles
  - Lines



High-level  
Objects:

- Lions
- Giraffes
- Zebras
- ...

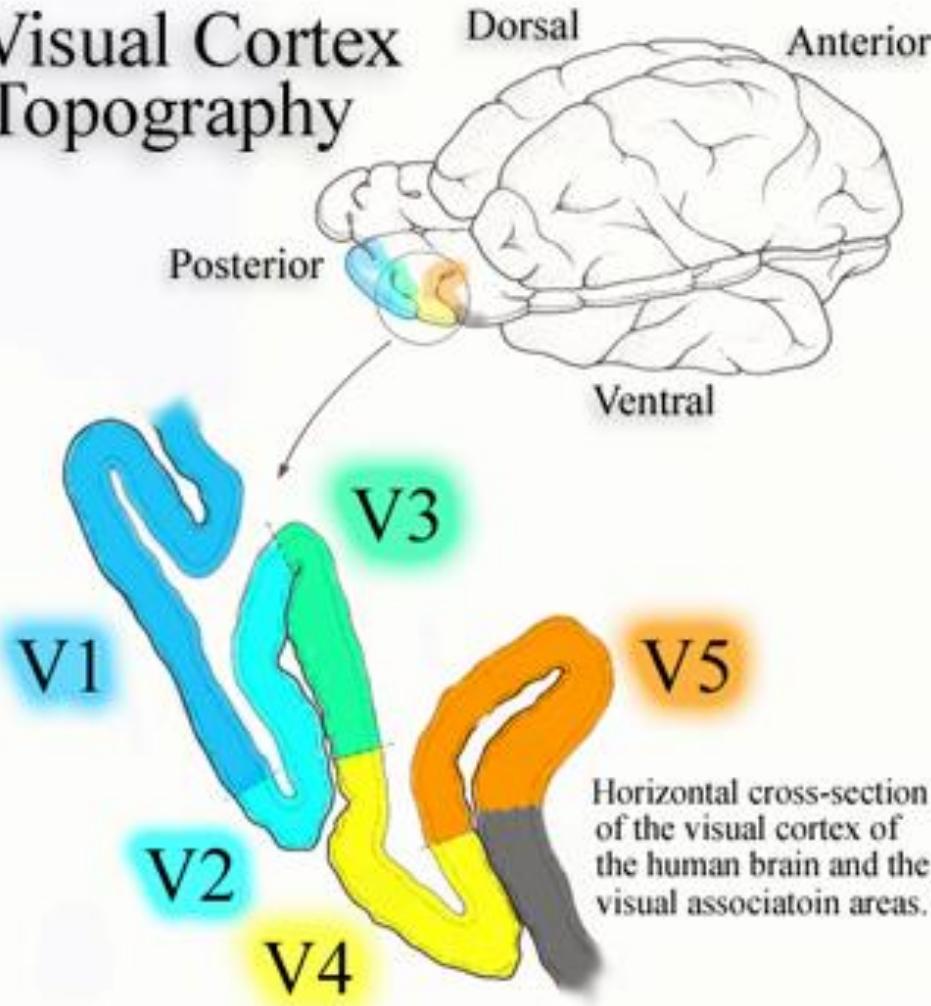


Object parts:

- Paws
- Legs
- Necks
- Heads
- ...

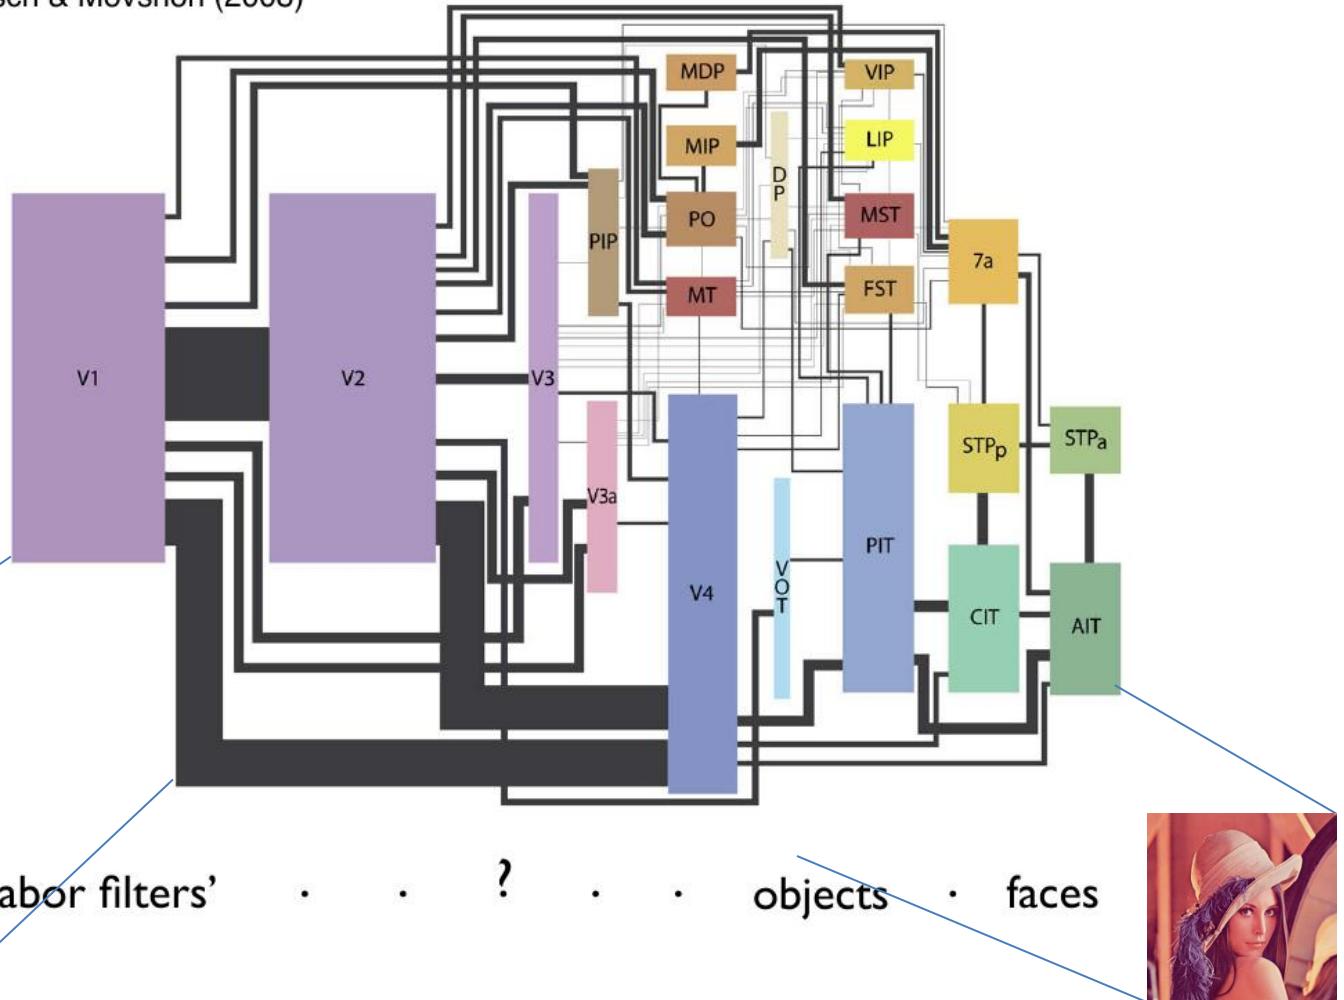
# Brain & Visual Cortex

## Visual Cortex Topography

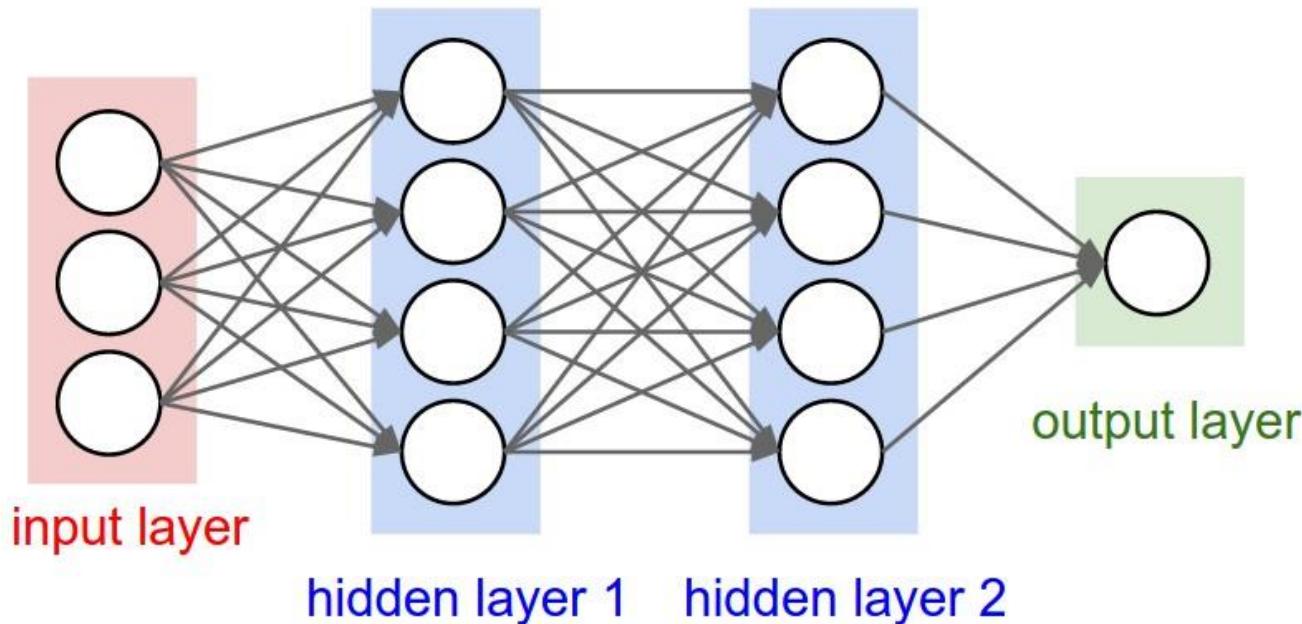


# Visual Cortex Diagram

Wallisch & Movshon (2008)



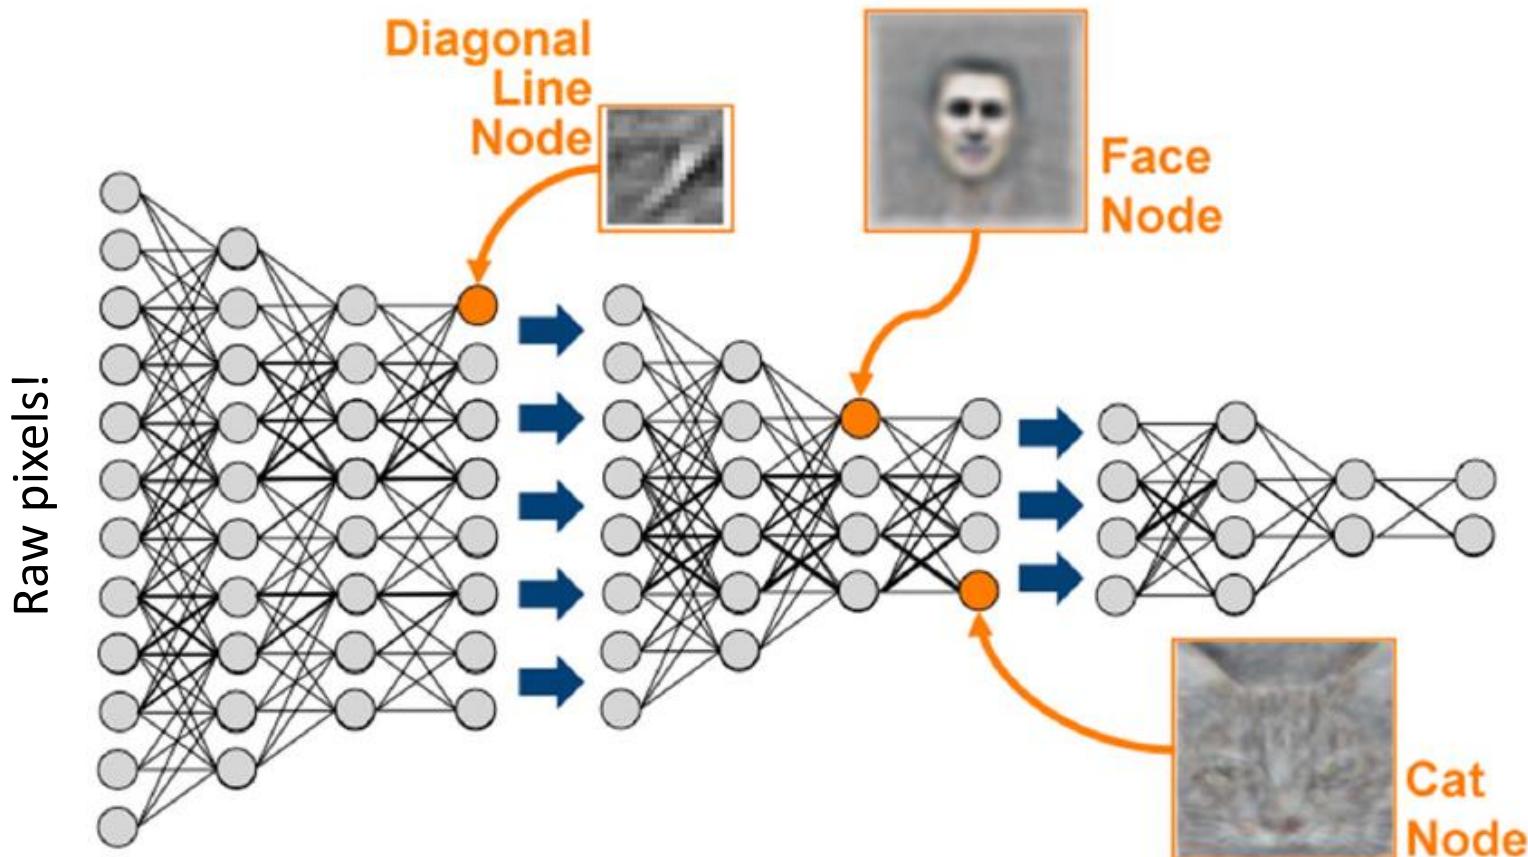
# Neural networks



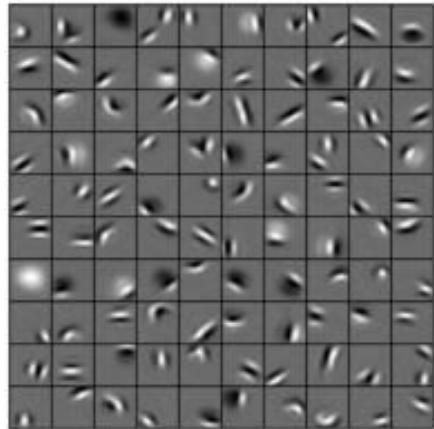
- A neuron detects some patterns in its inputs – combinations that cause it to fire
- When assembled into a network, neurons deep in the network react to patterns composed of more primitive parts

# Neural nets learn hierarchies!

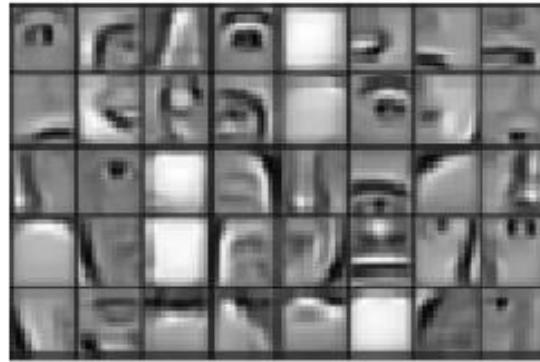
Google trained a network on YouTube videos. The net developed units detecting persons and cats!



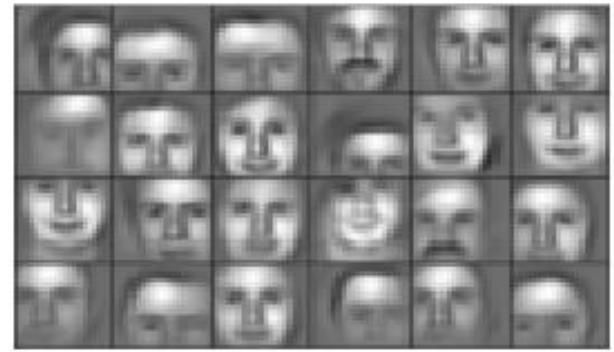
# Neural nets learn hierarchies!



First layer



Second layer



Third layer

Hierarchical features learned from a dataset of face images

(Lee et al., „Unsupervised Learning of Hierarchical Representations with Convolutional Deep Belief Networks“)

# Low-level features

What the neuron (feature-detector looks for)

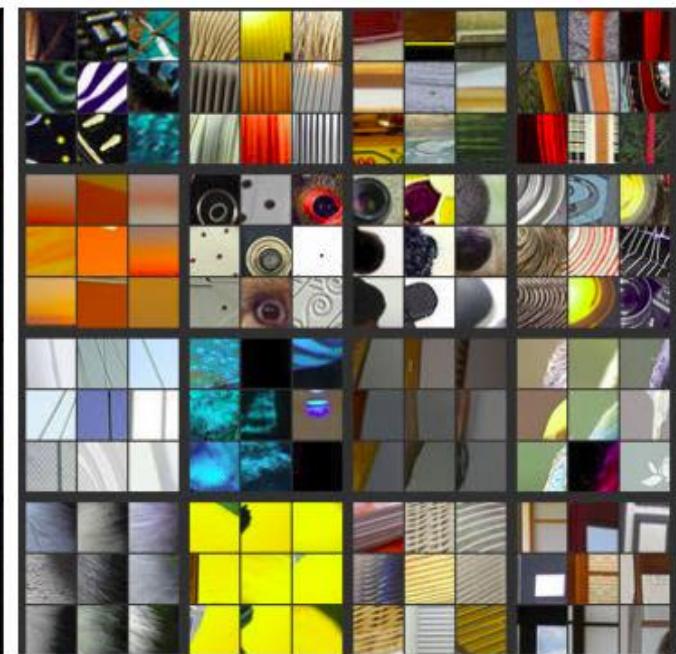
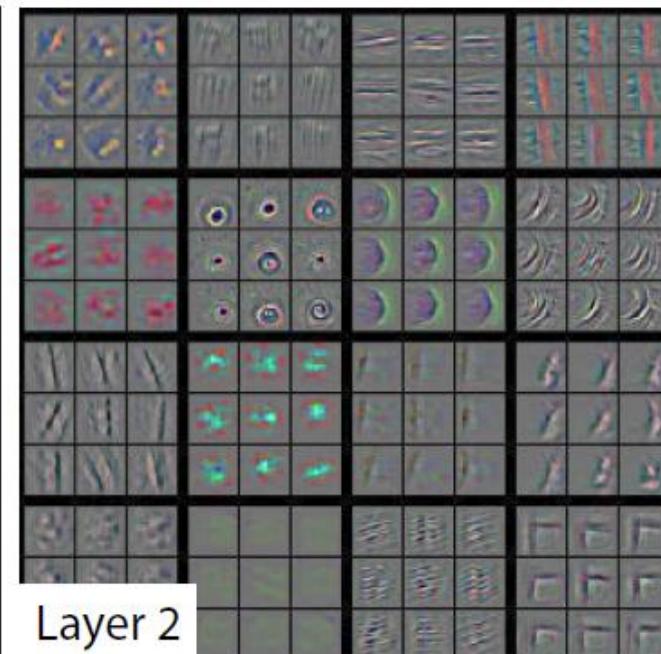


Layer 1

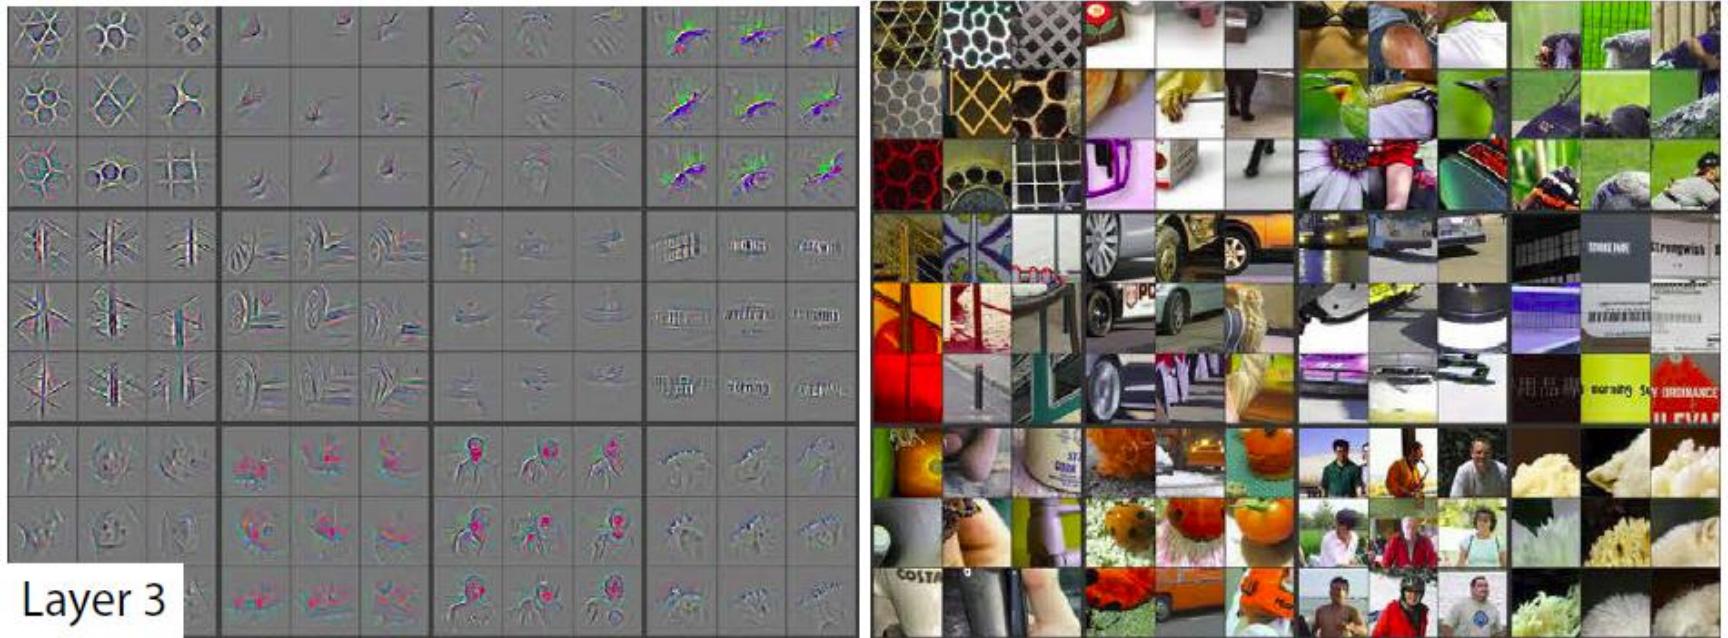


Layer 2

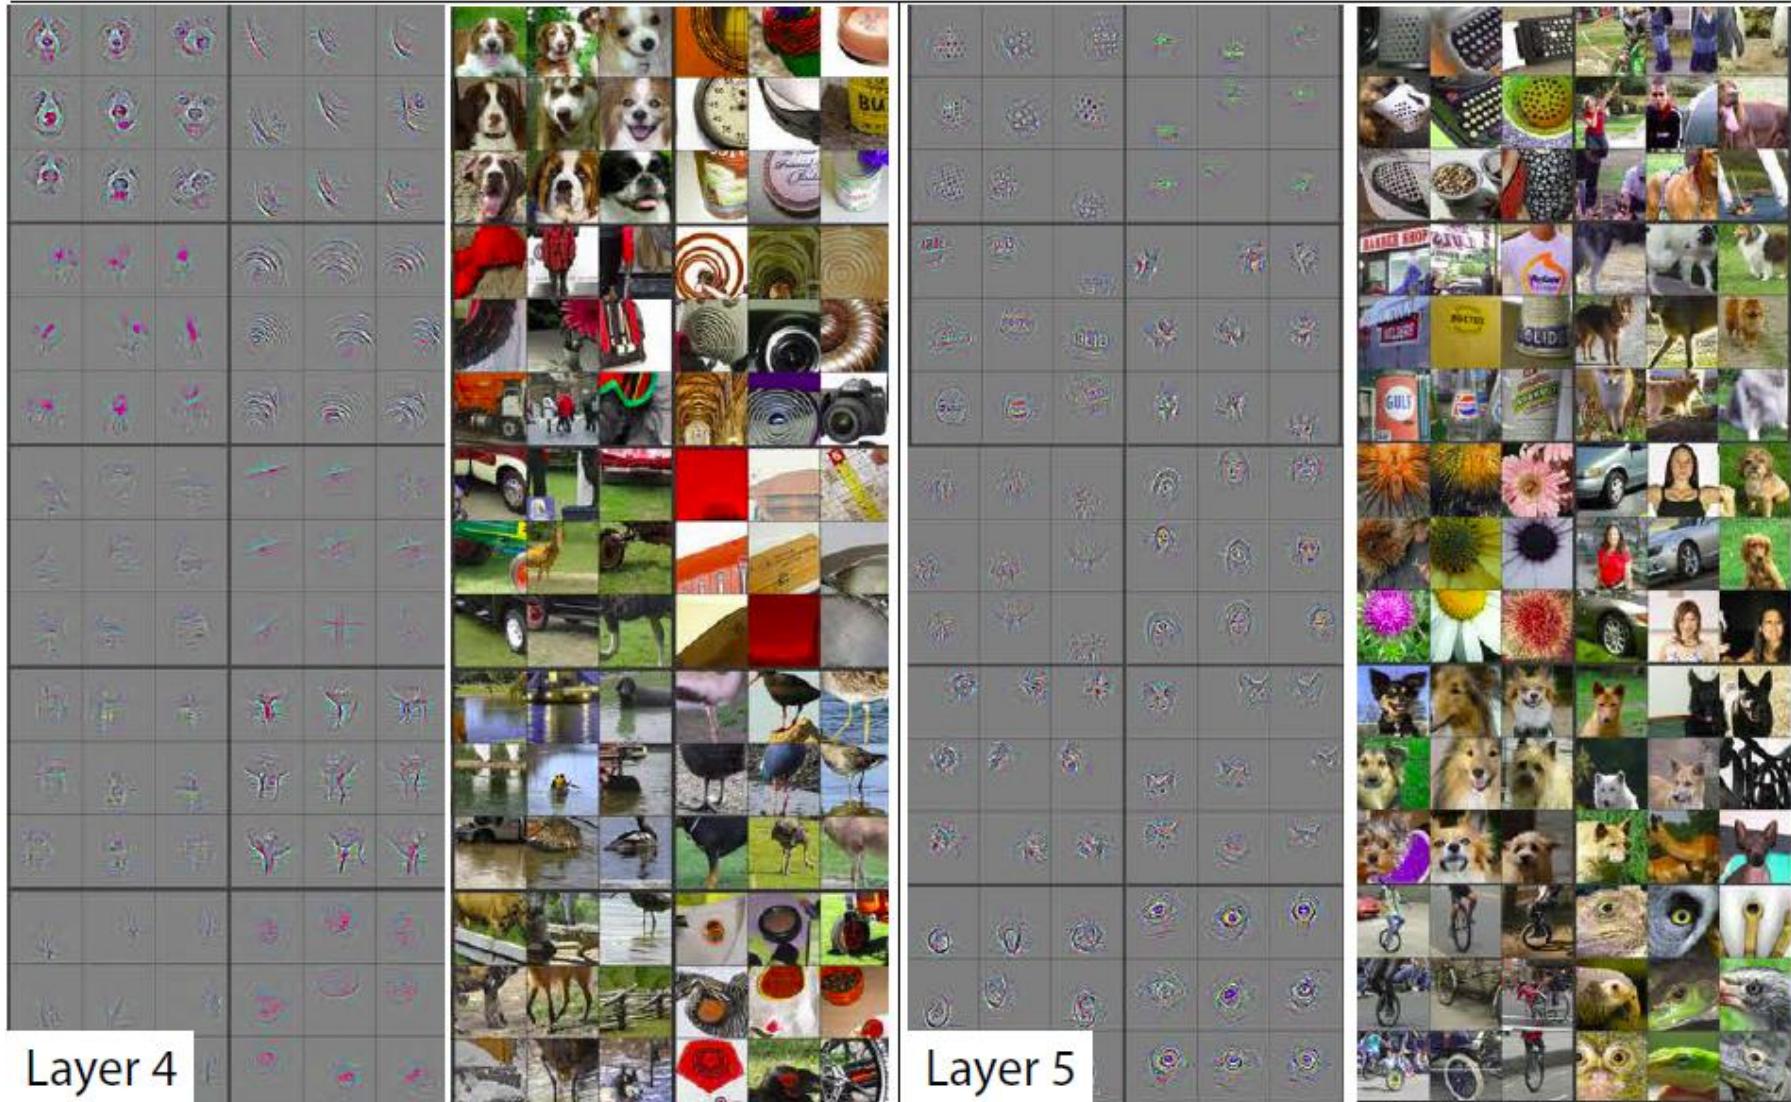
What images are selected by the neuron



# Mid-level features

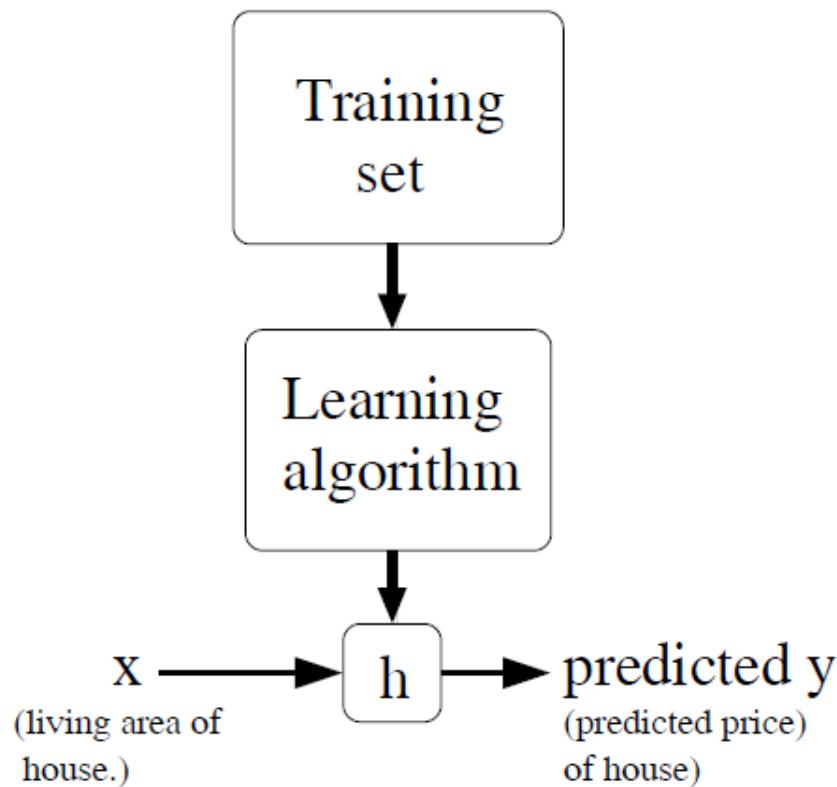


# High-level features



# Quick summary

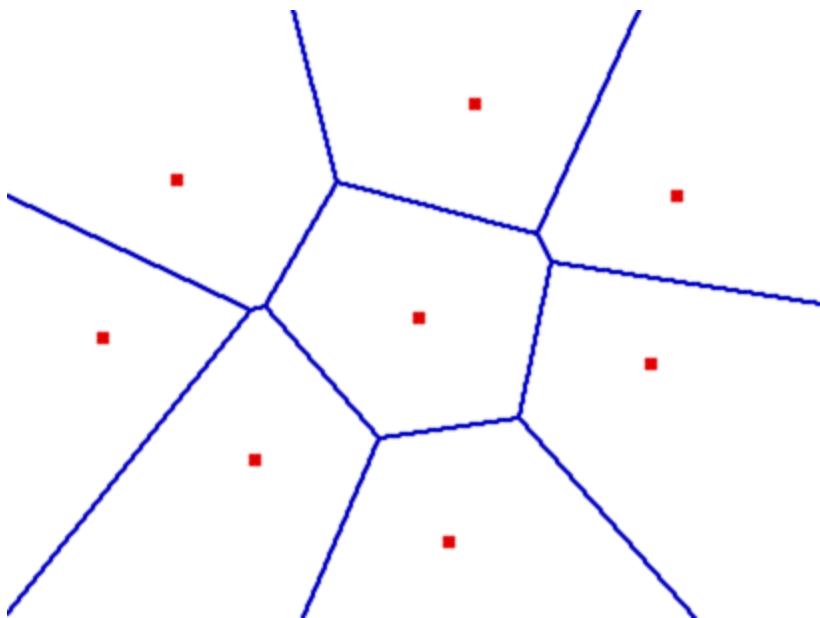
- ML algorithms distill data into models.
- We will call this learning from examples.



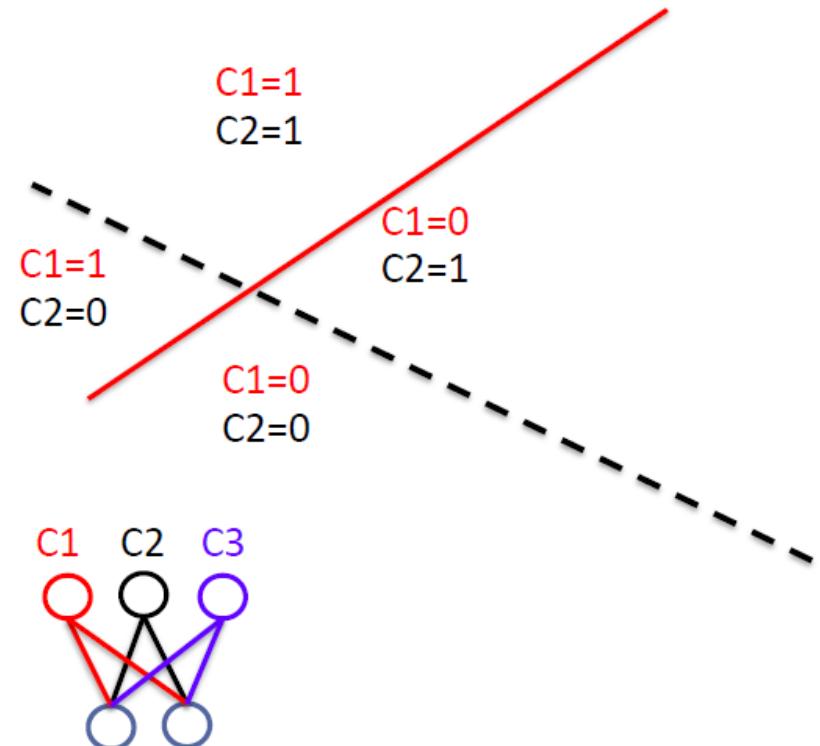
# Two kinds of models

## Nearest neighbors

- Look-up tables



## Neural nets



Bengio, 2009, Foundations and Trends in Machine Learning

# When to use machine learning

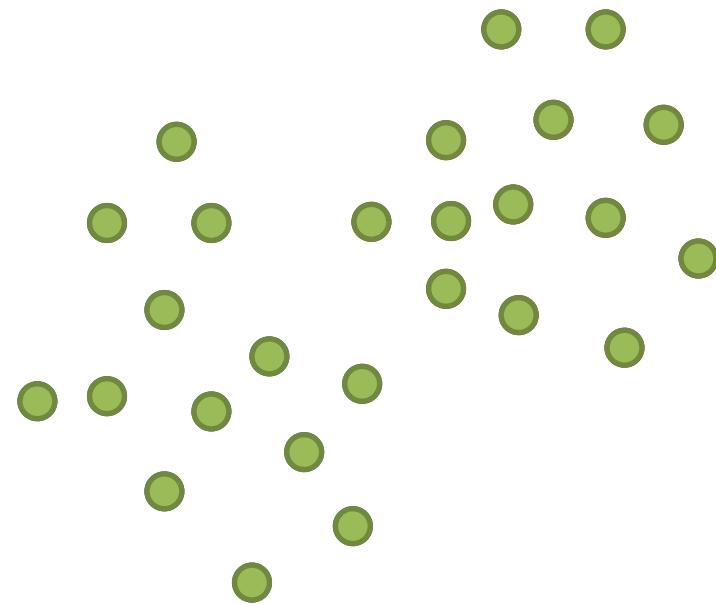
- Easy to get examples, hard to come up with an algorithm
- Don't know how to program a solution (e.g. speech recognition, language processing, translation rely heavily on data)
- We need to automatically tune or adapt the solution to the user
- The solution changes over time
- Question: when shouldn't we use learning?

# Programming vs. learning

- Requires thorough problem understanding
- Formally define pre- and post-conditions
- Implement the solution
- Prove the correctness
- Can have only a partial understanding of the problem – intuitions and a-priori assumptions
- Collect many examples of input-output
- Crucial aspect – **generalization** – will the learned solution work on new data?

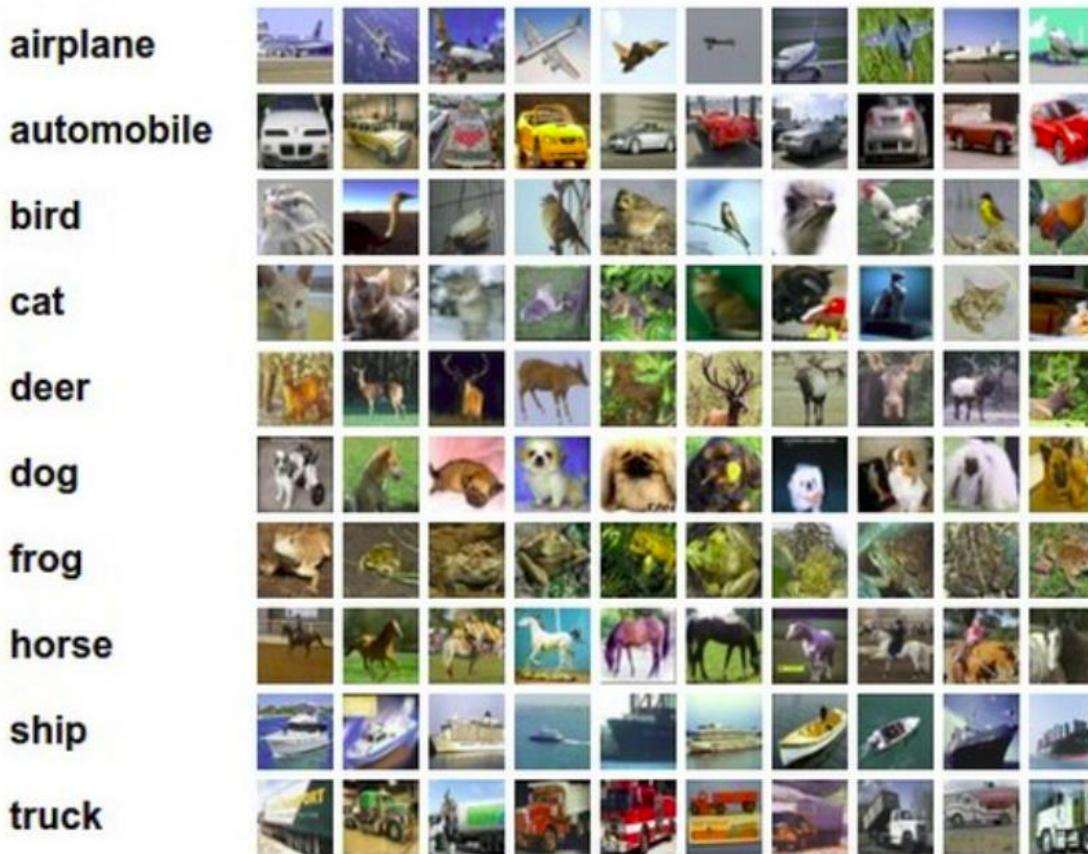
# Types of learning

- Supervised:
  - Learn an input-output relation, instant feedback
- Unsupervised:
  - Learn the density of data
  - Learn how to generate data
- Reinforcement:
  - Reward information given after a series of actions
  - Think of learning strategy in games



# Supervised learning examples

- Assign small images to one of ten categories





# “Life” Using Cucumbers

Launcher

How a Japanese cucumber farmer is using deep learning and TensorFlow

2L	
L	
M	
S	
2S	
BL	
BM	
BS	
C	



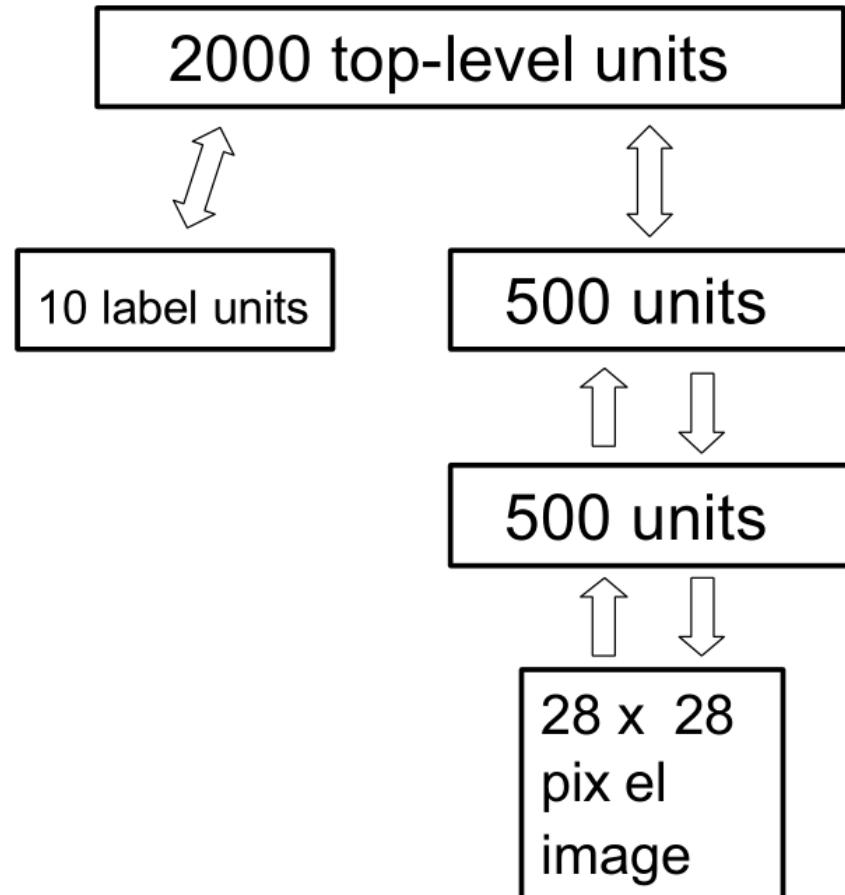
# Unsupervised learning example

- Generating album covers



Image by Alec Radford

# Digit Generation



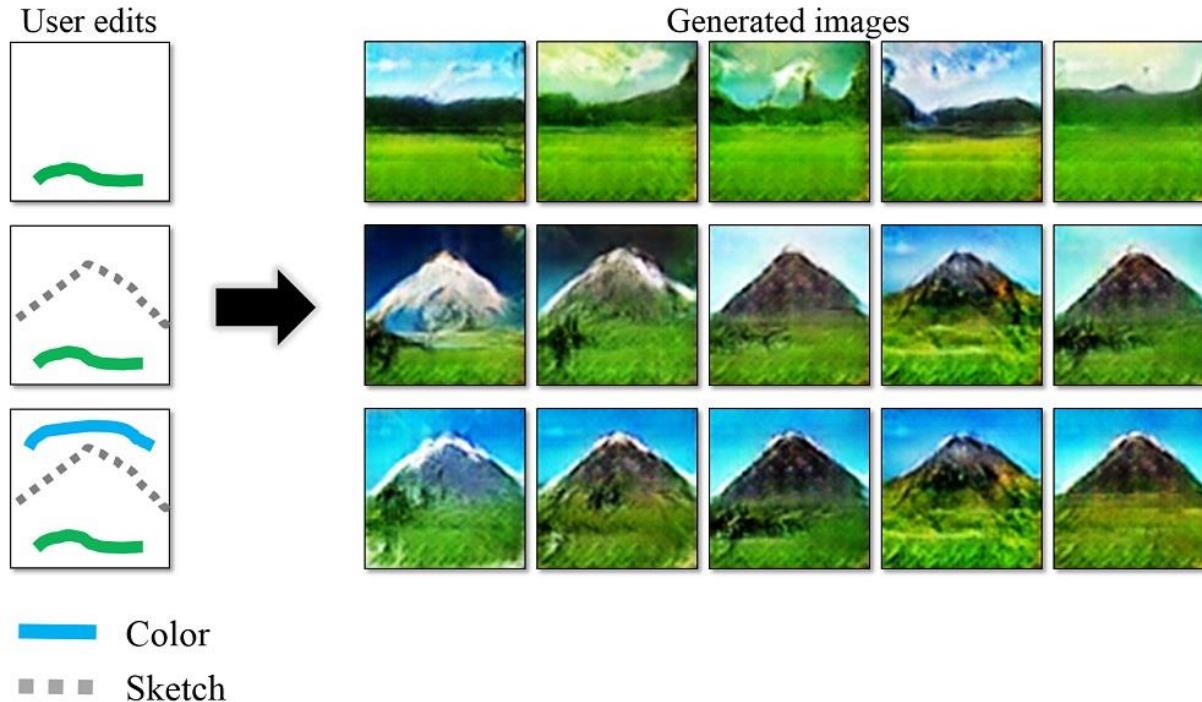
Demo: <http://www.cs.toronto.edu/~hinton/adi/index.htm>

# Brainstorming ideas

What would you do with a good generative model of images?

# Image manipulation

Transform sketches into images:



<https://www.youtube.com/watch?v=FDELBFSeqQs>

<https://www.youtube.com/watch?v=9c4z6YsBGQ0>

# Image super-resolution



(e) Bicubic



(f) SRCNN



(g) A+



(h) RAISR





+



=



# Style transfer

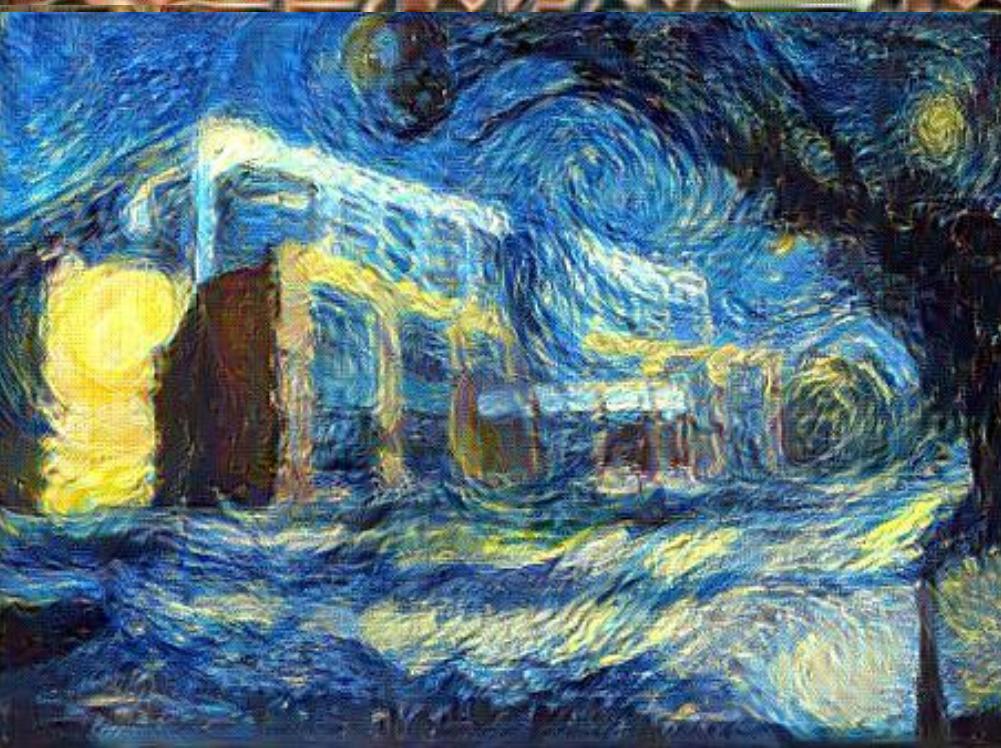
Find image that takes content from image A and style from B

Gatys et al., „A Neural Algorithm of Artistic Style”, 2015

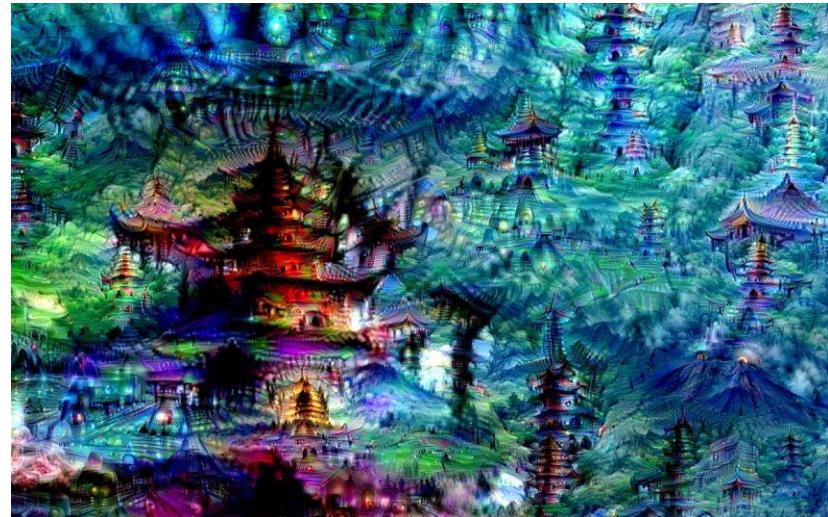
Sample adaptation to videos:

<https://www.youtube.com/watch?v=Khuj4ASldmU>

Ruder et al., „Artistic Style Transfer For Videos”



# Change the image to see many eyes/buildings in it.



Inceptionism: Going Deeper into Neural Networks

<http://googleresearch.blogspot.com/2015/06/inceptionism-going-deeper-into-neural.html>

Grocery Trip: <https://www.youtube.com/watch?v=DgPaCWJL7XI>

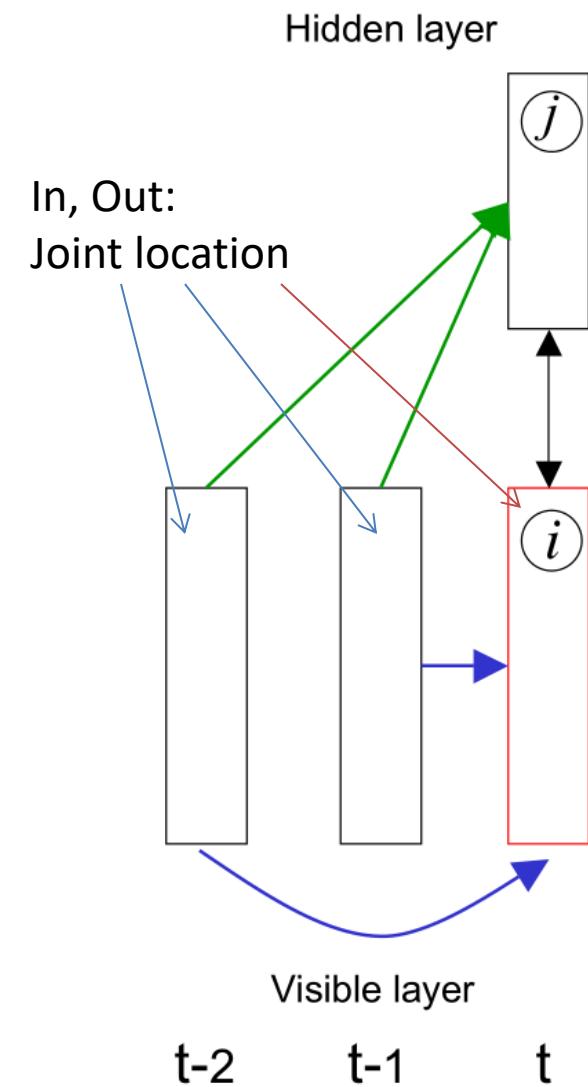
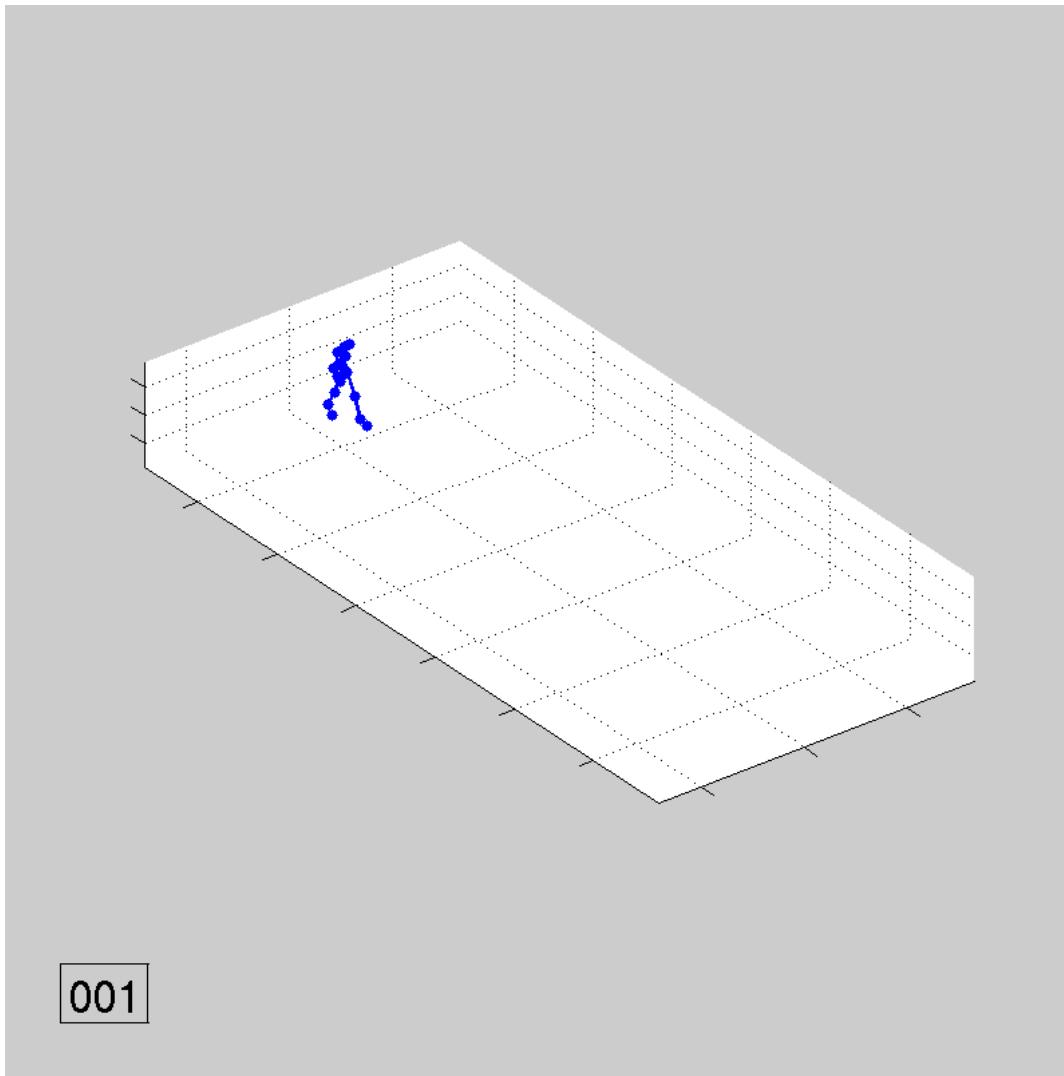
# In between supervised and unsupervised learning

- Sequence generation
  - Inherent targets -> predict the next/neighboring elements
- Typical use: language modeling

“Good everybody. Thank you very much. God bless the United States of America, and has already began with the world’s gathering their health insurance. It’s about hard-earned for our efforts that are not continued.”

(RNN generated Obama speech – by @samim)

# Network to generate skeleton animations



# Reinforcement learning

- Learning to play games
  - Historically: backgammon
  - Now: Atari games  
<https://www.youtube.com/watch?v=EfGD2qveGdQ>
- Learning to perform movements (robotics)  
<https://www.youtube.com/watch?v=EtMyH--vnU&t=25m>

# Supervised and Reinforcement Mix

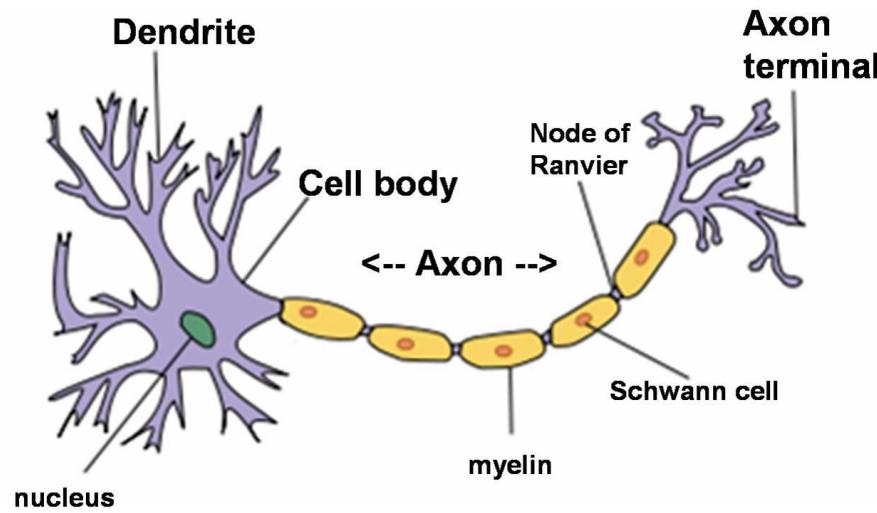
- AlphaGO:
  - 2 neural networks:
    - Position evaluator
    - Move proposer
  - First train to mimic human moves
  - Then do self-play and reinforcement learning

# Diving into Neural Networks

- All demos used Neural Networks – they are the best 😊
- Historical use: handwritten character recognition
- Nowadays: translation, image recognition, speech recognition, language processing...
- Current interest: large networks (Deep Learning), hierarchies of features,

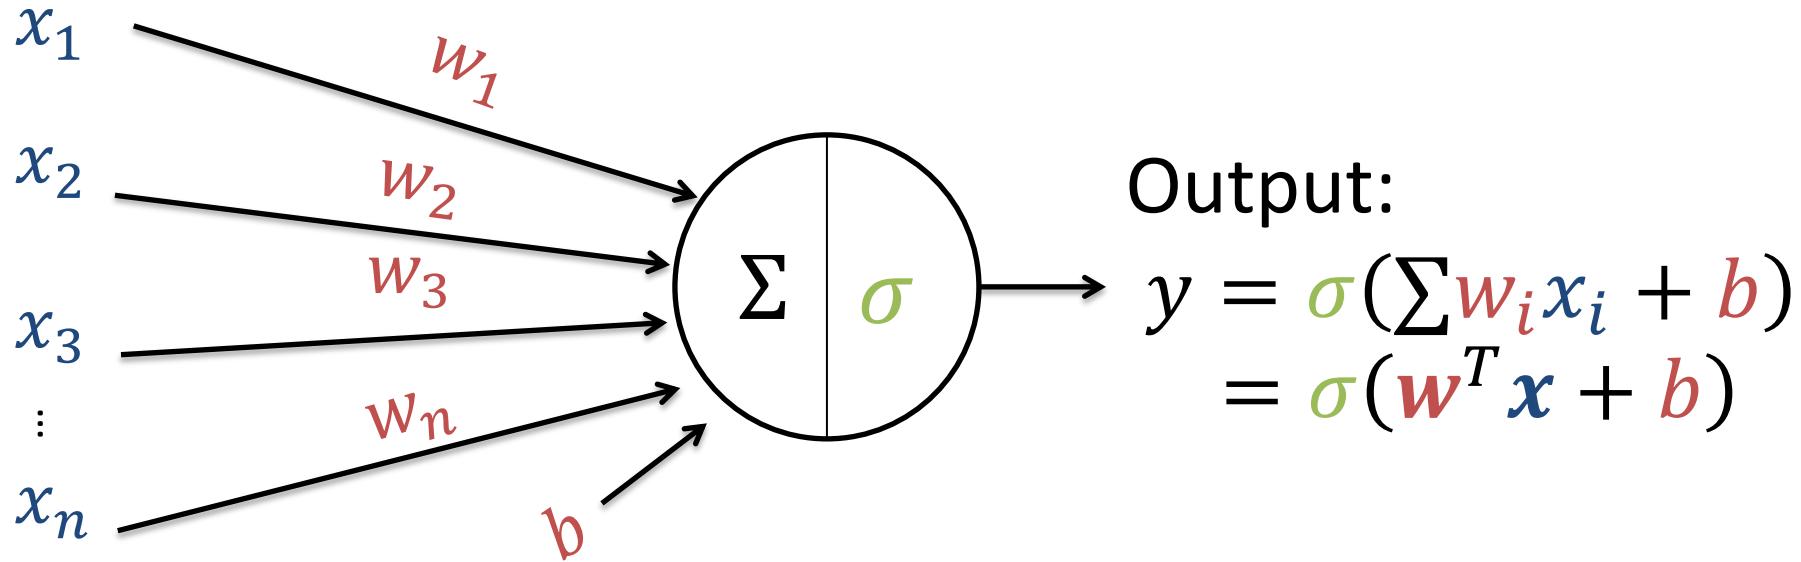
# Artificial Neural Networks

- Models remotely inspired by the brain cells
- A network of simple computational units resembling the biological neuron



- Attention: the brain is much more complicated (and very different) than ANNs!

# The artificial neuron (perceptron)

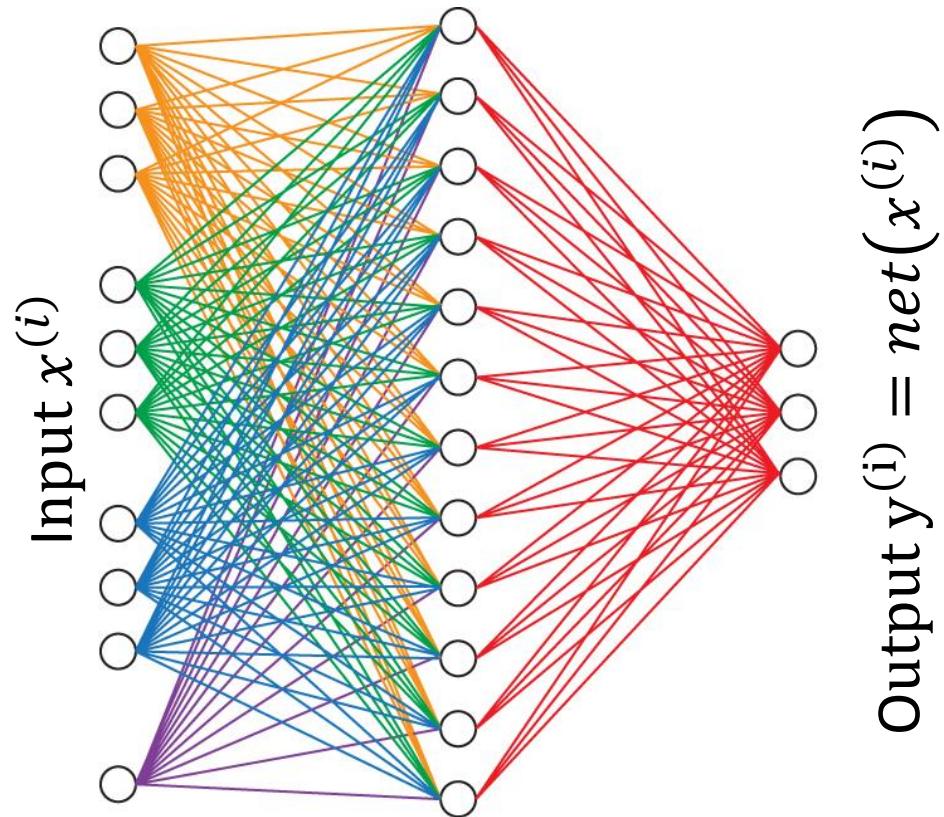


- $x_i$  are the inputs
  - $w_i$  are the weights and  $b$  the bias
  - $\Sigma$  denotes the summation
  - $\sigma$  is a (possibly nonlinear) activation function
- w<sub>i</sub>, b are TUNABLE!!*

# Neural (Perceptron) Network

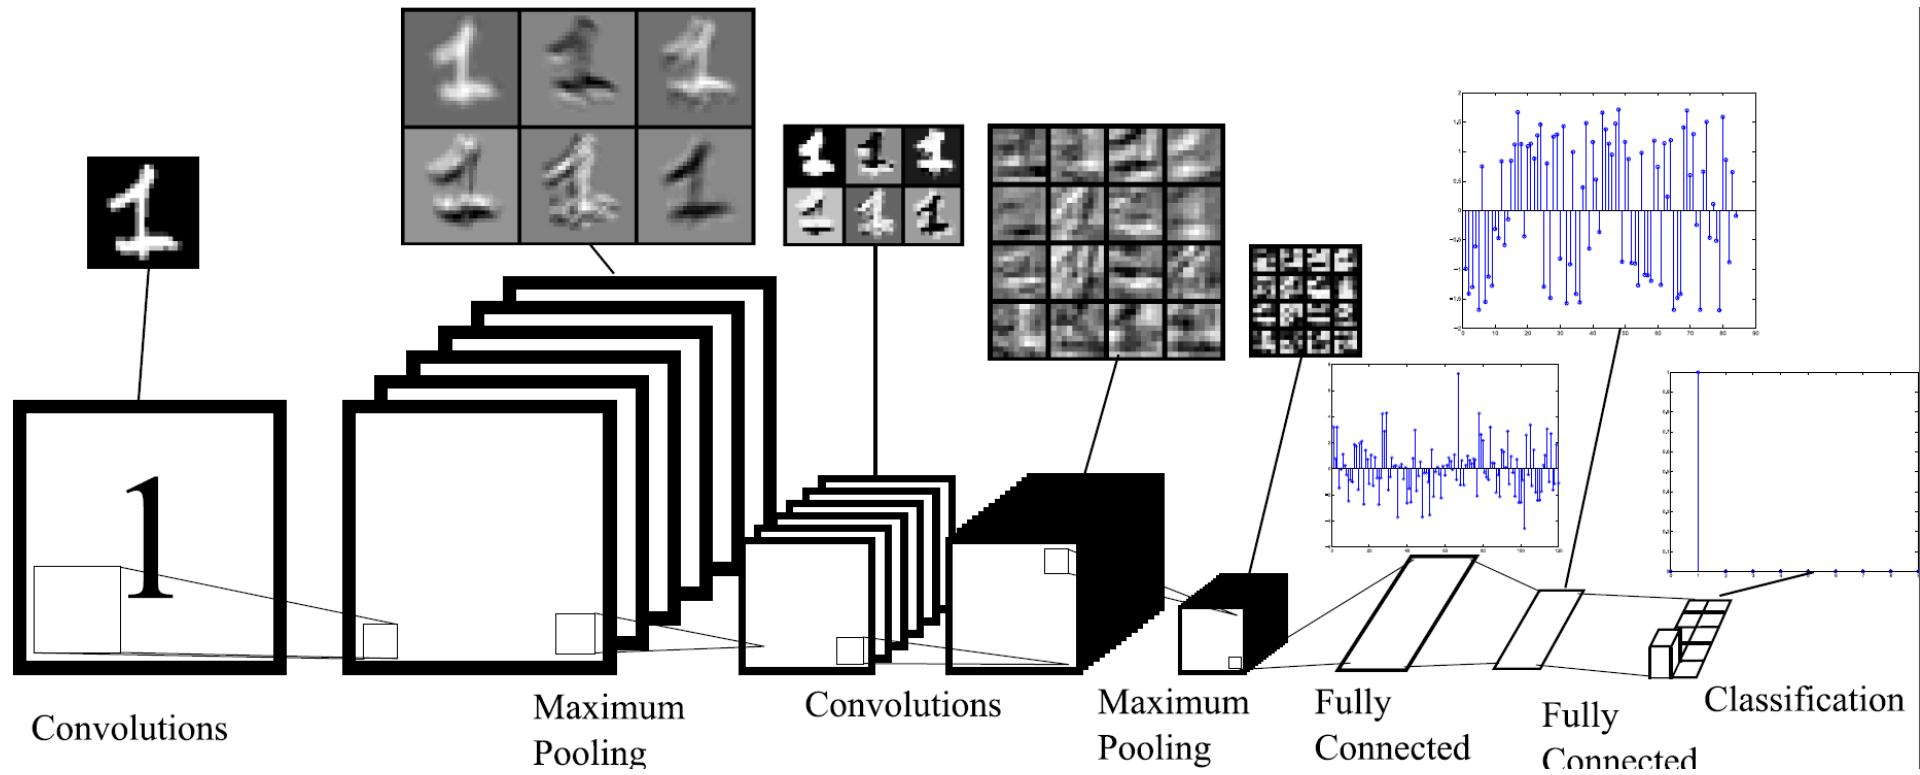
Data: pairs  $(x^{(i)}, y^{(i)})$

Learning:  $\min_{w,b} \sum_i (y^{(i)} - \text{net}(x^{(i)}))^2$



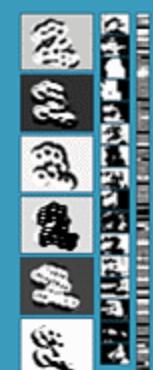
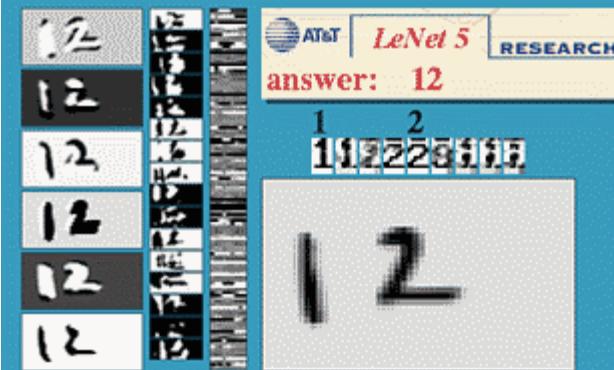
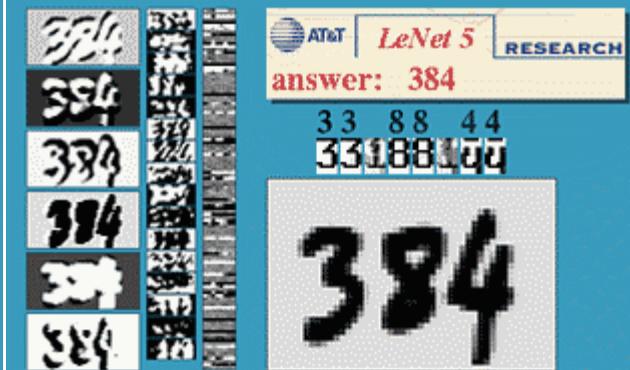
# Convolutional networks

Best for images!



# Convnets on digits

output



W5

W1

W3

input

# More demos

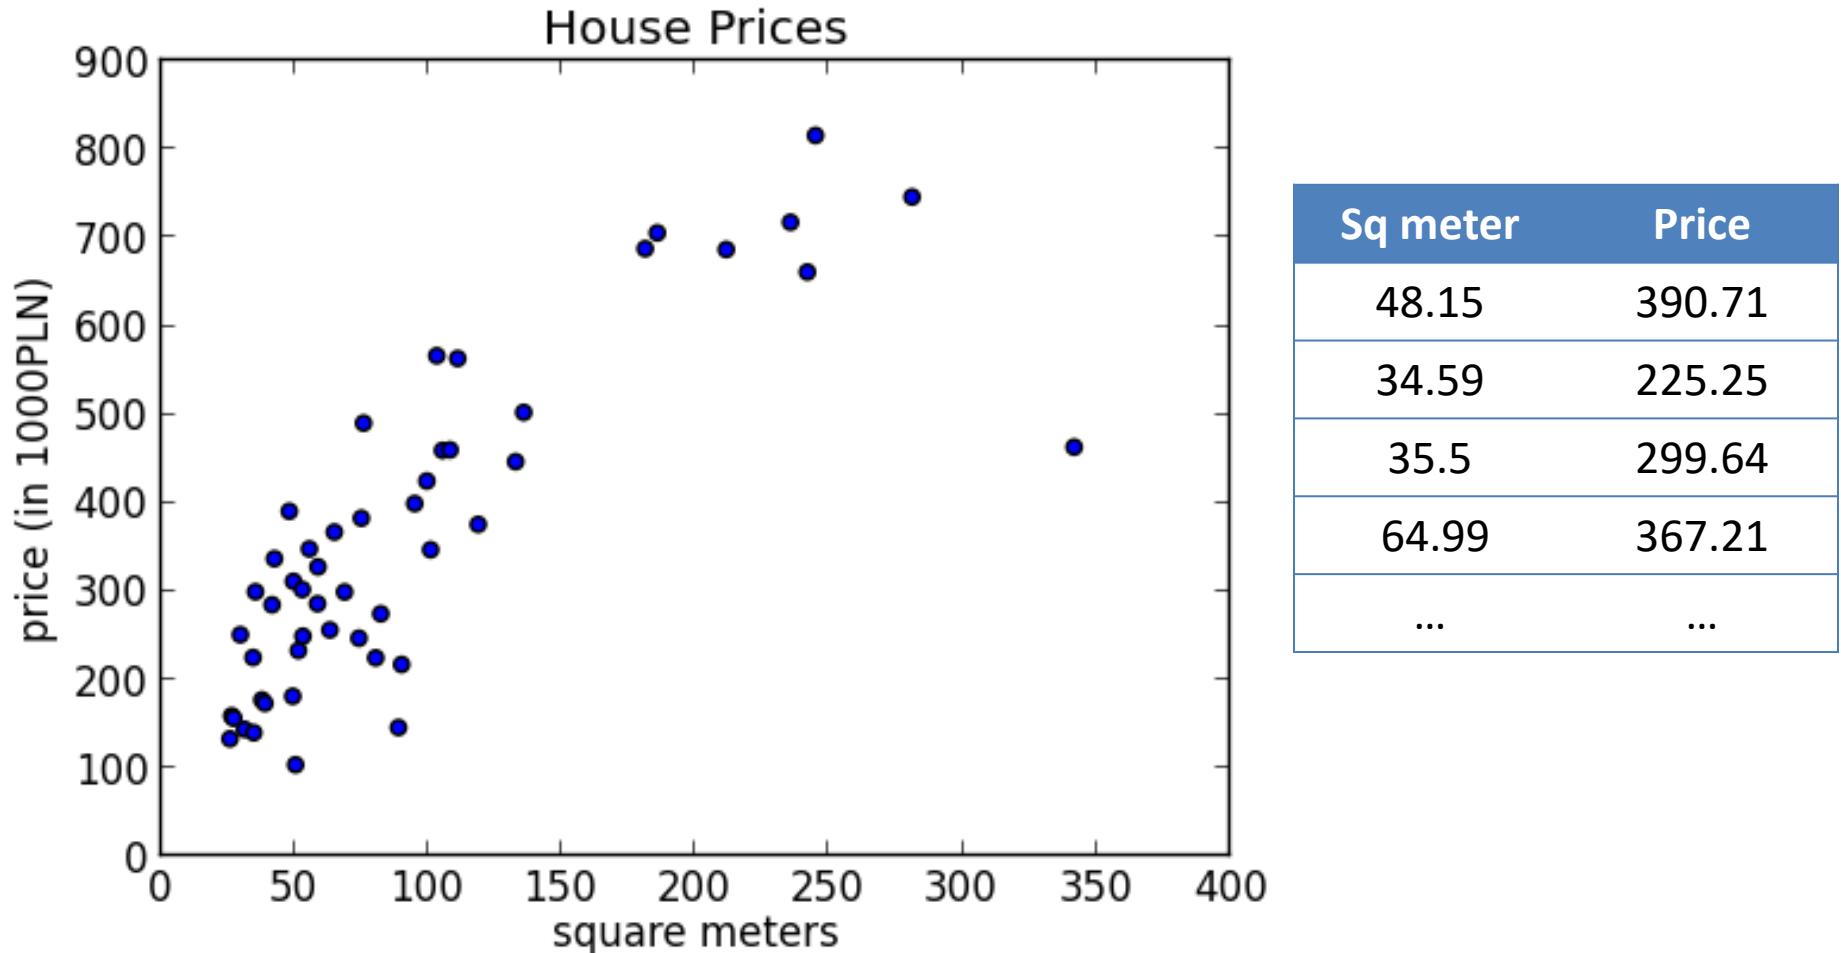
- [http://vdumoulin.github.io/morphing faces/online demo.html](http://vdumoulin.github.io/morphing_faces/online_demo.html)
- <http://www.cs.toronto.edu/~graves/handwriting.html>
- <http://deeplearning.net/demos/>
- <http://cs.stanford.edu/people/karpathy/convnetjs/>

# Intuitive description of learning

1. For a collected data set (training examples)
2. define a set of hypotheses that can describe the training set
3. then choose the best hypothesis based on some cost criterion and the data.

This course will be mainly about useful hypotheses and good cost functions, with a little bit of optimization.

# Example: predict prices



# Notation

- Vectors will be in **bold**
- $\mathbf{x}^{(i)}$  is the  $i$ -th training example
- $y^{(i)}$  is the  $i$ -th target (sometimes a vector too)
- $\{(\mathbf{x}^{(i)}, y^{(i)}) : i = 1 \dots m\}$  is the **training set**
- $h_{\Theta}$  will be the hypothesis
- $\Theta$  is the vector of  $h$ 's parameters (weights)
- $J(\Theta)$  will be the cost function

# Back to the example

- Out hypothesis:

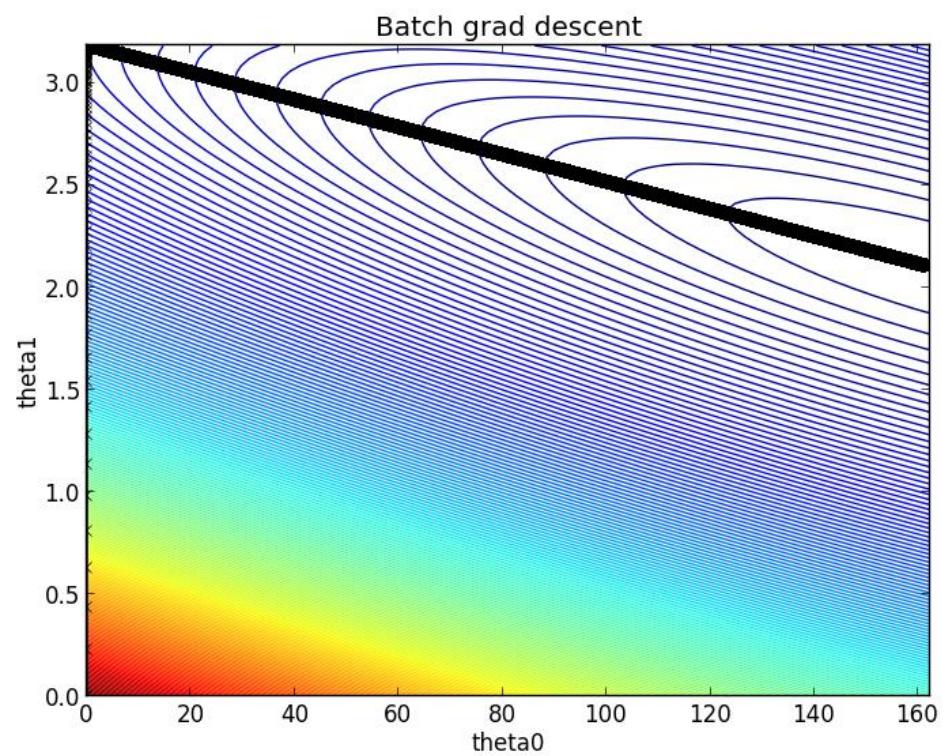
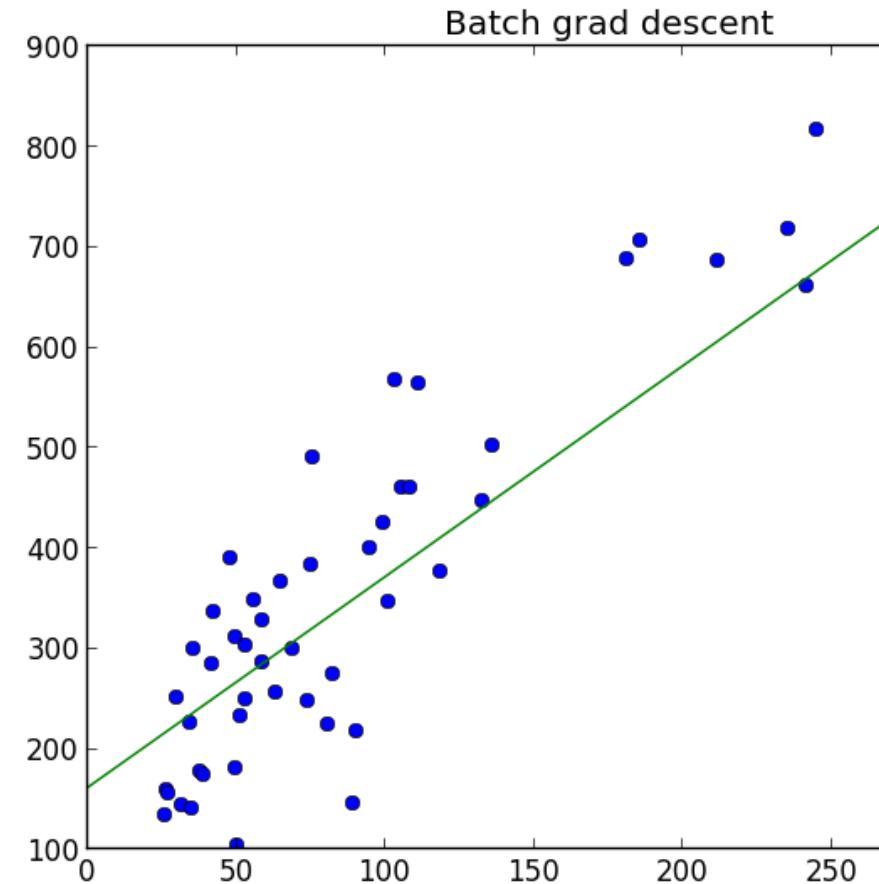
# Our cost function

- Measure how close are the predictions to the training data:

# How to choose $h$ ?

- Idea: pick one that minimizes  $J$
- Gradient descent – small steps towards the minimum
- Repeat:

# Example operation



# Neural network interpretation

# Gradient computation

# Analysis for a single example

# Two learning modes

- Batch gradient descent
  - Each gradient computation uses all examples
  - Exactly corresponds to minimization of  $J(\Theta)$
- On-line (stochastic) gradient descent
  - Each update computes the gradient on just **one** example or a few examples
  - Faster, but may not converge to the local minimum (will wander around)

# Batch/On-line comparison

