

Neural Networks and Deep Learning

Lecture 1: introduction

Jan Chorowski
Instytut Informatyki
Wydział Matematyki i Informatyki Uniwersytet
Wrocławski
2017-2018

Practical Information

- Website:
<http://ii.uni.wroc.pl/~jch/teaching/fall2018-2019/neural-networks-fall-18>
- Free on-line Textbooks:
 - <http://www.deeplearningbook.org/> -> You can also buy a printed version!
 - <http://neuralnetworksanddeeplearning.com/>
- Good lecture notes for 50% of material in this course:
<http://cs229.stanford.edu/>
- More advanced, very good courses on deep learning:
 - <https://ift6266h15.wordpress.com/>
 - <http://cs231n.stanford.edu/>
 - <http://cs224d.stanford.edu/>
- Good paper textbooks:
 - Goodfellow, Bengio, Courville „Deep Learning”
<http://www.deeplearningbook.org/>
 - Murphy, Machine Learning: a Probabilistic Perspective
<http://www.cs.ubc.ca/~murphyk/MLbook/>
 - Bishop, Pattern Recognition and Machine Learning (PRML)
<https://www.microsoft.com/en-us/research/people/cmbishop>

Why are you here?


What do you want to learn?

1. Machine Learning: using data to teach machines useful skills
2. Neural Networks: good Machine Learning models

Where is machine learning?

Everywhere!

- web search ☺
- and ads ☹
- recommendations
- Self-driving cars



Local results for **starbucks** near **Chicago, IL**

Ads

Starbucks Get Local Directi...
Phone Numbers |
MapQuest.com

Starbucks Ct Whatever you're l...
you can get it on
www.eBay.com

Buy Starbuck Find Starbucks C...
eBay Express Of...
www.eBayExpres...

Local Search Results

Starbucks in Chicagoland
This friendly neighborhood Starbucks is extra-spacious, ... of local hero Joe DiMaggio in this first Starbucks in the Little Italy neighborhood of Chicago. ...
www.starbuckseverywhere.net/Chicagoland.htm - 127k -
Cached - Similar pages

Starbucks in Illinois
Illinois Chicagoland - Illinois Remote
www.starbuckseverywhere.net/illinois
Cached - Similar pages

Google Organic Search Results

amazon.com

Hello, Scott Wheeler. We have recommendations for you. (Not Scott?)
Scott's Amazon.com Today's Deals Gifts & Wish Lists Gift Cards

Shop All Departments Search Amazon.com

Scott, Welcome to Your Amazon.com (If you're not Scott Wheeler, click here.)

Today's Recommendations For You


Here's a daily sample of items recommended for you. Click here to [see all recommendations](#).

Russia Map by ITMB International Travel Maps By ITMB Publishing Ltd \$10.95

In Search of Sunrise, Vol. 2: Asia by Ol' Theta ShinkinBank! \$13.99

Land of the Horizons: A History of the Oba by Jason Goodwin \$11.95


ML example: face detection



Source: Microsoft and Apple face detection API documentations, wykop.pl


Types of learning

- Supervised:
 - Learn an input-output relation, instant feedback
- Unsupervised:
 - Learn the density of data
 - Learn how to generate data
 - Group data into clusters
- Reinforcement:
 - Reward information given after a series of actions
 - Think of learning strategy in games



Supervised learning examples

- Assign small images to one of ten categories





“Life” Cucumbers

Launcher

How a Japanese Cucumber Farmer is Using Deep Learning and TensorFlow

2L	
L	
M	
S	
2S	
BL	
BM	
BS	
C	




Unsupervised learning example

- Generating album covers




Image by Alec Radford

Digit Generation



Demo: <http://www.cs.toronto.edu/~hinton/adi/index.htm>

ML Example: Google Photos app



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

Algos and Data

Sup. learning

Unsup. learning













Image recognition



Object Detection


<https://www.youtube.com/watch?v=WZmSMkK9VuA>






(Ren, He, Girshick and Sun, „Faster R-CNN”, 2015)

Natural Language Processing

<http://zapotoczny.pl/parser/>



NLP, embeddings



Using both language and images

Putting images and tags into the same vector space
(Kiros, Salakhutdinov, Zemel, TACL 2015)






- bowl + box =



- box + bowl =



Caption generation (Xu et al ICML 2015)




How are these problems solved?

1. Take **data** and a **learning algorithm**
2. The algorithm discovers **patterns** in the data and produces a **model**
3. Query model to ask questions about the data
 - „Is there any face in the image?”
 - “Is this review favorable?”
 - „What is the object in the image?”
 - “What caption is most likely given this image?”

Example: Digit Recognition

- Task: recognize handwritten digits
- Input: images 28×28 pixel values ($[0,1]^{784}$)
- Output: $\{0,1,\dots,9\}$




The easiest algorithm

- Collect examples – pairs of (input, output):

5	0	4	1	9	2	1	3	1	4
5	0	4	1	9	2	1	3	1	4

- To classify a new instance 3:
 - Find the most similar element of the train set 3
 - Return its label
- Seems too easy
 - It works well when we have enough data

KNN (k Nearest Neighbors)



Design decision: choose “ k ” – how many neighbors to use?
Few – sensitive to outliers, Many- very smooth classification boundary

KNN uses: recommendations

- Items are similar if users rate them similarly
- Training (offline):
 - Represent each item by its scores
 - Compute the distance between scores
- Recommendation generation (online):
 - Find nearest neighbors for each item recently browsed/put into the cart

Amazon patented this:




US007113917B2

(12) **United States Patent**
Jacobi et al.

(10) Patent No.: US 7,113,917 B2
(45) Date of Patent: *Sep. 26, 2006

Nearest Neighbors use: inpainting



Original Image

Input



Scene Matches

Output

1. Find similar photos
 - Use Gist descriptors
2. Match and copy

Works poorly with 1000 images...

But is good with
 >1000000

Alternative realities



Input



Alternative inpainting results

kNN for dialogs

A system Adrian Łaćucki coded based on Reddit.

Real (cherry-picked) conversation from Festiwal nauki:

User: hi

Bot: Hi.

User: i like you

Bot: Your in if u want.

User: lmao

Bot: Ayy lmao.

User: ayyy

Bot: Who are you?

User: im a nurse

Bot: Gott im Himmel. You have become the host.

User: and what are you ?

Bot: I am just a talking robot. But I like my job.

User: thats great


KNN – what makes images similar?



Quite similar



Not quite similar




Opposite




KNN – what makes images similar?

Pixel-wise difference can be large for visually similar images


Image




Rotated Image
diff=76.177262



Translated Image
diff=114.726490



Other Image
diff=56.507974




KNN requires lots of data

- Consider recognizing postal codes (5 digits)
- Will this work if we treat each postal code as a single image (10^5 combinations)?
- If we treat each digit separately, how to accommodate for:
 - Uniform style (slant, thickness) of digits in one code.
 - Segmentation, joins, overlaps of digits.

Curse of dimensionality

In highly dimensional spaces, things are dissimilar (far away from each other):

- Let N be the number of points.
- The average distance between points grows with data dimensionality.
- In other words, there are fewer neighbors within a radius from each point.



Humans see differently

Change blindness

- Humans are blind even to very drastic changes in images.
- We clearly treat things as the same even if the pixel-wise differences are huge!
- Example:

<https://www.youtube.com/watch?v=1nL5ulsWMYc>

<https://www.youtube.com/watch?v=FWSxSQsspiQ>

Learning is about invariants and generalizations



We want a hierarchical model



Low-level
concepts:

- Edges
- Points
- Gradients
-



Higher-level:

- Textures:
 - Spots
 - Stripes
- Geom.
shapes:
 - Corners
 - Circles
 - Lines



High-level
Objects:

- Lions
- Giraffes
- Zebras
- ...




Object parts:

- Paws
- Legs
- Necks
- Heads
- ...


Brain & Visual Cortex

Visual Cortex Topography




Visual Cortex Diagram

Wallisch & Movshon (2008)




Neural networks




- A neuron detects some patterns in its inputs – combinations that cause it to fire
- When assembled into a network, neurons deep in the network react to patterns composed of more primitive parts

Neural nets learn hierarchies!


Google trained a network on YouTube videos. The net developed units detecting persons and cats!



Neural nets learn hierarchies!



First layer



Second layer



Third layer

Hierarchical features learned from a dataset of face images

(Lee et al., „Unsupervised Learning of Hierarchical Representations with Convolutional Deep Belief Networks”)

Low-level features

What the neuron (feature-detector looks for)





Layer 1




Layer 2


What images are selected by the neuron



Mid-level features




High-level features



Quick summary


- ML algorithms distill data into models.
- We will call this learning from examples.




Two kinds of models

Nearest neighbors

- Look-up tables



Neural nets



Bengio, 2009, Foundations and Trends in Machine Learning

When to use machine learning

- Easy to get examples, hard to come up with an algorithm
- Don't know how to program a solution (e.g. speech recognition, language processing, translation rely heavily on data)
- We need to automatically tune or adapt the solution to the user
- The solution changes over time
- Question: when shouldn't we use learning?

Programming vs. learning


- Requires thorough problem understanding
- Formally define pre- and post-conditions
- Implement the solution
- Prove the correctness
- Can have only a partial understanding of the problem – intuitions and a-priori assumptions
- Collect many examples of input-output
- Crucial aspect – **generalization** – will the learned solution work on new data?

Brainstorming ideas

What would you do with a good generative model of images?

Image manipulation

Transform sketches into images:



<https://www.youtube.com/watch?v=FDELBFSeqQs>

<https://www.youtube.com/watch?v=9c4z6YsBGQ0>

Image super-resolution



(e) Bicubic



(f) SRCNN



(g) A+




(h) RAISR





+



=



Style transfer

Find image that takes content from image A and style from B

Gatys et al., „A Neural Algorithm of Artistic Style”, 2015



Sample adaptation to videos:

<https://www.youtube.com/watch?v=Khuj4ASldmU>

Ruder et al., „Artistic Style Transfer For Videos”



Change the image to see many eyes/buildings in it.



Inceptionism: Going Deeper into Neural Networks

<http://googleresearch.blogspot.com/2015/06/inceptionism-going-deeper-into-neural.html>

Grocery Trip: <https://www.youtube.com/watch?v=DgPaCWJL7XI>



In between supervised and unsupervised learning

- Sequence generation
 - Inherent targets -> predict the next/neighboring elements
- Typical use: language modeling

“Good everybody. Thank you very much. God bless the United States of America, and has already began with the world’s gathering their health insurance. It’s about hard-earned for our efforts that are not continued.”

(RNN generated Obama speech – by @samim)

Network to generate skeleton animations



Reinforcement learning

- Learning to play games
 - Historically: backgammon
 - Now: Atari games
<https://www.youtube.com/watch?v=EfGD2qveGdQ>
- Learning to perform movements (robotics)
<https://www.youtube.com/watch?v=EtMyH--vnU&t=25m>

Supervised and Reinforcement Mix

- AlphaGO:
 - 2 neural networks:
 - Position evaluator
 - Move proposer
 - First train to mimic human moves
 - Then do self-play and reinforcement learning